Additional Curated Provider Candidates — Round 2 Live-Probe Research (2026-06-13)

Round-2 research feeding **Defect B** (on-device api-app embeds curated Go providers in `lava-api-go/internal/provider/curated/`).

Already shipped (6) — NOT re-proposed here: The Pirate Bay (`apibay.org`), YTS (`yts.mx`), Torrents-CSV (`torrents-csv.com`), BitSearch (`bitsearch.eu`), Knaben (`api.knaben.org`), Nyaa (`nyaa.si` RSS).

Already rejected with evidence (round 1 + prior) — NOT re-probed: EZTV (keyword no-op), 1337x (CF challenge), TorrentGalaxy (DNS dead), SolidTorrents (= BitSearch backend), Torlock (HTML timeout), torrent-api-py (404).

All probes were live, server-side `curl` GETs from a macOS host on 2026-06-13 (~09:38–11:50 UTC), bounded ≤18s each, default UA `Mozilla/5.0 (compatible; LavaResearchBot/1.0)` unless noted. Per §11.4.6 no-guessing vocabulary: only captured facts are stated; anything not directly probed is marked `UNCONFIRMED:`.

The bar (ALL must hold, each PROVEN with a captured probe)

1. Free-text query genuinely FILTERS — two distinct queries return DIFFERENT result sets that each contain the query term (the EZTV no-op test).
2. Anonymous — no API key, no login, no cookie.
3. NOT Cloudflare-challenge-gated for a plain server-side Go GET (no HTTP 403 + `cf-mitigated: challenge` / "Just a moment").
4. Returns an `info_hash` or magnet.
5. Stable JSON (preferred) or RSS / trivially parseable response.

Candidate results

1. TorrentDownloads (torrentdownloads.pro) — VIABLE (TOP PICK)

- **Search endpoint:** GET https://www.torrentdownloads.pro/rss.xml?type=search&search=<query> (a www.torrentdownloads.pro host; the bare host issues a redirect — follow it / pin www.).
- **Live probe:** HTTP 200, body is RSS XML (<?xml ... <rss version='2.0'>), bytes≈26627 for q=ubuntu. server: cloudflare present but **cf-mitigated ABSENT** — passive CDN, NOT a challenge gate. Works with a plain non-browser UA (LavaServerBot/1.0 → 50 items), so no UA-gating.
- **Auth:** anonymous; no key/cookie.
- **info_hash / magnet:** YES — dedicated <info_hash> element, 40-hex lowercase (e.g. 19987310611f0fae3b2c9678601b92967d938aff). Magnet is built client-side from the hash + the curated public-tracker list (same pattern as the Nyaa provider).

2-query filter-difference proof (the EZTV no-op test) — PASSES decisively:

```
Q1 search=ubuntu → 50 items; 50/50 <title> contain "ubuntu"
   ("Ubuntu 12 10 Server 64 Bit", "Ubuntu 10 10 i386 v14 2 5", ...)
Q2 search=debian → 50 items; 50/50 <title> contain "debian"
overlap between the two title sets: 0
```

Field mapping (RSS <item> → provider.SearchItem):

RSS element	provider.SearchItem	Notes
<title>	Title	plain text
<info_hash>	InfoHash	40-hex lowercase; build magnet from it
<size>	SizeBytes	integer bytes (e.g. 24898149)
<seeders>	Seeders	integer (e.g. 978)
<leechers>	Leechers	integer (e.g. 30)
<pubDate>	Date	RFC-1123 (Mon, 21 Oct 2013 13:20:49 +0200)
<link>	(detail URL)	site-relative /torrent/<id>/<slug>

VERDICT: VIABLE. Best round-2 candidate — full field set (the ONLY round-2 candidate with native <info_hash> + <seeders> + <leechers> + <size> all in the feed), clean filter, passive CF, no UA-gating. General-purpose corpus (complements the niche providers). Maps cleanly onto the nyaa RSS provider shape. **Caveat (UNCONFIRMED:):** .pro is one of several rotating TorrentDownloads mirror TLDs; long-term host stability of .pro specifically is not proven beyond today's probe.

---

2. Tokyo Toshokan (tokyotosho.info) — VIABLE (anime/niche, no seeders)

- **Search endpoint:** GET https://www.tokyotosho.info/rss.php?terms=<query>
- **Live probe:** HTTP 200, content-type: application/rss+xml; charset=utf-8. server: cloudflare, cf-mitigated ABSENT — passive CDN.
- **Auth:** anonymous.
- **info_hash / magnet:** YES — each item's <description> (CDATA) embeds a full magnet:?xt=urn:btih:<BASE32-40>&tr=... link (e.g. magnet:?xt=urn:btih:AIHBWH5HWF2UQCHMDT2D3O6TP63HBLNX&tr=...). Note the btih is **base32** (32 chars), not 40-hex — the provider must accept the magnet as-is OR base32-decode → hex.

2-query filter-difference proof — PASSES:

```
Q1 terms=naruto → 139 items; first content titles on-topic
   ("Cutie Kuromi - Naruto chan.zip", "...(NARUTO -ナル
ト-).zip", ...)
Q2 terms=bleach → 146 items; on-topic
   ("Bleach The new Gotei-13 law - Final [ENG].zip", "...
(Bleach)...", ...)
disjoint, query-specific result sets.
```

Field mapping (RSS <item> → provider.SearchItem):

RSS element	provider.SearchItem	Notes
<title>	Title	plain text
magnet: URL inside <description>	InfoHash / magnet	base32 btih; regex-extract from CDATA
Size: NNN MB inside <description>	SizeBytes	human-readable; must parse "113.89MB"
<pubDate>	Date	RFC-1123 GMT
<category>	(category)	e.g. "Other", "Hentai (Manga)"
seeders / leechers	— (ABSENT)	**NOT in the feed →** emit 0/unknown

VERDICT: VIABLE. Honest CapSearch (genuine filter), anonymous, magnet present. **Two caveats:** (a) NO seeders/leechers in the feed (must surface as 0/unknown — UI honesty); (b) btih is base32, and size needs human-string parsing. Niche (anime/Japanese media) — overlaps Nyaa's domain but draws a distinct corpus + different uploaders. Lower priority than TorrentDownloads because of the missing seeders + base32 + the adult-heavy default result mix.

---

3. LimeTorrents (`limetorrents.lol` → `.fun`) — VIABLE (redirect + indirect hash)

- **Search endpoint:** `GET https://www.limetorrents.lol/searchrss/<query>/` — issues `HTTP 301 → https://www.limetorrents.fun/searchrss/<query>/` (final host serves the RSS). **Must follow the redirect** (the `.lol`/`.fun` TLD rotates; pinning a single TLD is fragile).
- **Live probe (final host):** `HTTP 200`, RSS XML, `server: cloudflare`, `cf-mitigated` ABSENT — passive CDN.
- **Auth:** anonymous.
- **info_hash / magnet:** INDIRECT — no `<info_hash>` element and no magnet; the 40-hex hash is embedded in the `<enclosure url>` pointing to `https://itorrents.net/torrent/<40HEX>.torrent` (e.g. `itorrents.net/torrent/232CD67EB3FFBD7C37BF9EC3EE887417E5AE1EE6.torrent`). Must regex-extract the `[A-F0-9]{40}` and build the magnet client-side.

2-query filter-difference proof — PASSES:

```
Q1 searchrss/ubuntu/ → 40 items; 41 <title> lines incl. ubuntu (41
= 40 items + feed header)
Q2 searchrss/debian/ → 40 items; 41 <title> lines incl. debian
overlap between the two title sets: 0
```

Field mapping (RSS `<item>` → provider.SearchItem):

RSS element	provider.SearchItem	Notes
`<title>`	`Title`	plain text
`<enclosure url>` (itorrents)	`InfoHash`	extract `[A-F0-9]{40}` from URL; build magnet
`<size>`	`SizeBytes`	integer bytes
`Seeds: N , Leechers M` in `<description>`	`Seeders/Leechers`	regex-parse the description string
`<pubDate>`	`Date`	e.g. `12 Dec 2025 10:15:20 +0200` (non-

RSS element	provider.SearchItem	Notes
		standard, no weekday)

VERDICT: VIABLE (lowest of the three). Genuine filter, passive CF, hash + seeders recoverable — but two friction points lower its rank: (a) hard dependency on a `.lol → .fun` redirect on a rotating-TLD host (the most fragile of the three); (b) hash is indirect (URL-embedded, uppercase hex) and seeders/leechers need free-text regex out of `<description>`. Acceptable as a later addition, behind TorrentDownloads + TokyoTosho.

REJECTED candidates (each with captured evidence)

GloDLS (`glodls.to → gtorrents.site`) — REJECTED (EZTV-class no-op)

- `GET https://glodls.to/rss.php?q=<query> → 302 → gtorrents.site/rss.php`, final `HTTP 200`, `application/xhtml+xml`.
- **Filter is a NO-OP** (the canonical EZTV trap): `q=ubuntu` and `q=debian` returned the **byte-identical 50-item "latest torrents" feed** — `<guid>` lists diff'd IDENTICAL; 0/50 titles contained the query term in either feed (first item for both was "Iron Maiden - Virtual XI ... [FLAC]"). The `q` parameter is ignored. **Fails bar #1.**

Torrent-Paradise (`.ml / .org / .eu`) — REJECTED (dead / parked)

- `torrent-paradise.ml/api/search?q=ubuntu → HTTP 200` but `content-type: text/html` serving a **Hungarian online-casino spam page** ("Magyar online kaszinók 2026 ...") — the domain has been repurposed; the JSON API is gone. `torrent-paradise.org` and `torrent-paradise.eu → curl (6) Could not resolve host`. **Fails bar #1/#5 (no API).**

AcademicTorrents (`academictorrents.com`) — REJECTED (no per-query search feed)

- The site explicitly **blocks browse-scraping** (`browse.php → 307 → checkb.htm`, an HTML page that says: "*Hello if you are reading this you may be trying to scrape the browse page ... we instead ask you to search an XML file (in RSS format)*"). The offered feeds — `rss.xml` (recent) and `database.xml` (full DB dump) — take **NO query parameter**: `rss.xml` returns a fixed

44 KB "Recent Torrents" list regardless of input. `apiv2/.../search` → `404`. There is no anonymous free-text-filtering endpoint. **Fails bar #1.**

Anidex (`anidex.info`) — REJECTED (unstable / down)

- `GET https://anidex.info/rss/?q=naruto` → first attempt `curl (28)` timeout (15 s, 0 bytes); retry with a browser UA → `HTTP 502` (Cloudflare error page "Error 502"). The origin is unreachable/unstable. **Fails bar #3/#5.**

ext.to (`ext.to`, ExtraTorrent mirror) — REJECTED (403)

- `GET https://ext.to/rss/?q=ubuntu` → `HTTP 403` (`text/html`). Server-side GET is refused. **Fails bar #3.**

iDope (`idope.se`) — REJECTED (dead)

- `GET https://idope.se/torrent-list/ubuntu/?r=1` → `curl (28)` connection timeout (12 s). Host unreachable. **Fails bar #5.**

Snowfl (`snowfl.com`) — REJECTED (token-gated, not anonymous-stable)

- `GET https://snowfl.com/index1.php?ubuntu&<token>&...` → `HTTP 404` "Error". Snowfl's search requires a per-session token minted by its JS bundle; there is no stable anonymous query endpoint a plain Go GET can hit. **Fails bar #2/#5.**

Bitmagnet public instances — REJECTED (no stable public host)

- `bitmagnet.io/graphql` → `404`; `demo.bitmagnet.io` → DNS does not resolve. Bitmagnet is self-host software, not a public aggregator with a stable anonymous endpoint to pin. **Fails bar #5.**

BTDigg (`btdig.com`) — REJECTED (unreachable)

- `GET https://btdig.com/search?q=ubuntu&api=1` → `curl (28)` timeout (12 s). **Fails bar #5.**

Ranked VIABLE shortlist (ready to implement)

Rank	Provider	Endpoint	info_hash	Seeders/ Leechers
1	**TorrentDownloads**	`GET …torrentdownloads.pro/rss.xml?type=search&search=<q>`	native `<info_hash>` 40-hex	YES (native `<seeders>`/ `<leechers>`)
2	**Tokyo Toshokan**	`GET …tokyotosho.info/rss.php?terms=<q>`	magnet (base32 btih) in `<description>`	NO (absent)
3	**LimeTorrents**	`GET …limetorrents.lol/searchrss/<q>/` (→ `.fun`)	indirect (40-hex in itorrents enclosure URL)	YES (regex from `<description`

All three: anonymous, passive-Cloudflare (no `cf-mitigated` challenge), RSS, implementable against the `internal/provider/curated/nyaa` RSS pattern.

Single best next-to-implement

TorrentDownloads (`torrentdownloads.pro`). Strongest evidence line:

> `search=ubuntu` and `search=debian` each return 50 RSS items, 50/50 titles on-topic, **0 overlap**, with a native `<info_hash>` 40-hex element plus `<seeders>`/`<leechers>`/`<size>`/`<pubDate>` — and `server: cloudflare` with `cf-mitigated` ABSENT, served identically to a plain `LavaServerBot/1.0` UA (no browser/UA gating).

It is the only round-2 candidate that ships a complete native field set (hash + seeders + leechers + size + date) in the feed, making it the cleanest drop-in clone of the existing Nyaa provider.