

Additional Curated Provider Candidates — Round 2 Live-Probe Research (2026-06-13)

Round-2 research feeding **Defect B** (on-device api-app embeds curated Go providers in `lava-api-go/internal/provider/curated/`).

Already shipped (6) — NOT re-proposed here: The Pirate Bay (`apibay.org`), YTS (`yts.mx`), Torrents-CSV (`torrents-csv.com`), BitSearch (`bitsearch.eu`), Knaben (`api.knaben.org`), Nyaa (`nyaa.si` RSS).

Already rejected with evidence (round 1 + prior) — NOT re-probed: EZTV (keyword no-op), 1337x (CF challenge), TorrentGalaxy (DNS dead), SolidTorrents (= BitSearch backend), Torlock (HTML timeout), torrent-api-py (404).

All probes were live, server-side curl GETs from a macOS host on 2026-06-13 (~09:38–11:50 UTC), bounded ≤18s each, default UA Mozilla/5.0 (compatible; LavaResearchBot/1.0) unless noted. Per §11.4.6 no-guessing vocabulary: only captured facts are stated; anything not directly probed is marked UNCONFIRMED:.

The bar (ALL must hold, each PROVEN with a captured probe)

1. Free-text query genuinely FILTERS — two distinct queries return DIFFERENT result sets that each contain the query term (the EZTV no-op test).
 2. Anonymous — no API key, no login, no cookie.
 3. NOT Cloudflare-challenge-gated for a plain server-side Go GET (no HTTP 403 + `cf-mitigated: challenge` / "Just a moment").
 4. Returns an `info_hash` or magnet.
 5. Stable JSON (preferred) or RSS / trivially parseable response.
-

Candidate results

1. TorrentDownloads (torrentdownloads.pro) — VIABLE (TOP PICK)

- **Search endpoint:** GET `https://www.torrentdownloads.pro/rss.xml?type=search&search=<query>` (a `www.torrentdownloads.pro` host; the bare host issues a redirect — follow it / pin `www.`).
- **Live probe:** HTTP 200, body is RSS XML (`<?xml ... <rss version='2.0'>`), bytes≈26627 for `q=ubuntu`. server: `cloudflare` present but **cf-mitigated ABSENT** — passive CDN, NOT a challenge gate. Works with a plain non-browser UA (`LavaServerBot/1.0` → 50 items), so no UA-gating.
- **Auth:** anonymous; no key/cookie.
- **info_hash / magnet:** YES — dedicated `<info_hash>` element, 40-hex lowercase (e.g. `19987310611f0fae3b2c9678601b92967d938aff`). Magnet is built client-side from the hash + the curated public-tracker list (same pattern as the Nyaa provider).

2-query filter-difference proof (the EZTV no-op test) — PASSES decisively:

Q1 `search=ubuntu` → 50 items; 50/50 `<title>` contain "ubuntu"
("Ubuntu 12 10 Server 64 Bit", "Ubuntu 10 10 i386 v14 2 5", ...)
Q2 `search=debian` → 50 items; 50/50 `<title>` contain "debian"
overlap between the two title sets: 0

Field mapping (RSS `<item>` → `provider.SearchItem`):

RSS element	<code>provider.SearchItem</code>	Notes
<code><title></code>	Title	plain text
<code><info_hash></code>	InfoHash	40-hex lowercase; build magnet from it
<code><size></code>	SizeBytes	integer bytes (e.g. 24898149)
<code><seeders></code>	Seeders	integer (e.g. 978)
<code><leechers></code>	Leechers	integer (e.g. 30)
<code><pubDate></code>	Date	RFC-1123 (Mon, 21 Oct 2013 13:20:49 +0200)
<code><link></code>	(detail URL)	site-relative <code>/torrent/<id>/<slug></code>

VERDICT: VIABLE. Best round-2 candidate — full field set (the ONLY round-2 candidate with native <info_hash> + <seeders> + <leechers> + <size> all in the feed), clean filter, passive CF, no UA-gating. General-purpose corpus (complements the niche providers). Maps cleanly onto the nyaa RSS provider shape. **Caveat (UNCONFIRMED):** .pro is one of several rotating TorrentDownloads mirror TLDs; long-term host stability of .pro specifically is not proven beyond today's probe.

2. Tokyo Toshokan (tokyotosho.info) — VIABLE (anime/niche, no seeders)

- **Search endpoint:** GET <https://www.tokyotosho.info/rss.php?terms=<query>>
- **Live probe:** HTTP 200, content-type: application/rss+xml; charset=utf-8. server: cloudflare, cf-mitigated ABSENT — passive CDN.
- **Auth:** anonymous.
- **info_hash / magnet:** YES — each item's <description> (CDATA) embeds a full magnet:?xt=urn:btih:<BASE32-40>&tr=... link (e.g. magnet:?xt=urn:btih:AIHBWH5HWF2UQCHMDT2D306TP63HBLNX&tr=...). Note the btih is **base32** (32 chars), not 40-hex — the provider must accept the magnet as-is OR base32-decode → hex.

2-query filter-difference proof — PASSES:

Q1 terms=naruto → 139 items; first content titles on-topic
("Cutie Kuromi - Naruto chan.zip", "... (NARUTO -ナルト-).zip", ...)
Q2 terms=bleach → 146 items; on-topic
("Bleach The new Gotei-13 law - Final [ENG].zip", "... (Bleach)...", ...)
disjoint, query-specific result sets.

Field mapping (RSS <item> → provider.SearchItem):

RSS element	provider.SearchItem	Notes
<title>	Title	plain text
magnet: URL inside <description>	InfoHash / magnet	base32 btih; regex-extract from CDATA
Size: NNN MB inside <description>	SizeBytes	human-readable; must parse "113.89MB"
<pubDate>	Date	RFC-1123 GMT
<category>	(category)	e.g. "Other", "Hentai (Manga)"
seeders / leechers	— (ABSENT)	

RSS element	provider.SearchItem	Notes
		NOT in the feed → emit 0/unknown

VERDICT: VIABLE. Honest CapSearch (genuine filter), anonymous, magnet present. **Two caveats:** (a) NO seeders/leechers in the feed (must surface as 0/unknown — UI honesty); (b) btih is base32, and size needs human-string parsing. Niche (anime/Japanese media) — overlaps Nyaa's domain but draws a distinct corpus + different uploaders. Lower priority than TorrentDownloads because of the missing seeders + base32 + the adult-heavy default result mix.

3. LimeTorrents (limetorrents.lol → .fun) — VIABLE ⚠ (redirect + indirect hash)

- **Search endpoint:** GET <https://www.limetorrents.lol/searchrss/<query>/> — issues HTTP 301 → <https://www.limetorrents.fun/searchrss/<query>/> (final host serves the RSS). **Must follow the redirect** (the .lol/.fun TLD rotates; pinning a single TLD is fragile).
- **Live probe (final host):** HTTP 200, RSS XML, server: cloudflare, cf-mitigated ABSENT — passive CDN.
- **Auth:** anonymous.
- **info_hash / magnet:** INDIRECT — no <info_hash> element and no magnet; the 40-hex hash is embedded in the <enclosure url> pointing to <https://itorrents.net/torrent/<40HEX>.torrent> (e.g. itorrents.net/torrent/232CD67EB3FFBD7C37BF9EC3EE887417E5AE1EE6.torrent). Must regex-extract the [A-F0-9]{40} and build the magnet client-side.

2-query filter-difference proof — PASSES:

Q1 searchrss/ubuntu/ → 40 items; 41 <title> lines incl. ubuntu (41 = 40 items + feed header)

Q2 searchrss/debian/ → 40 items; 41 <title> lines incl. debian
overlap between the two title sets: 0

Field mapping (RSS <item> → provider.SearchItem):

RSS element	provider.SearchItem	Notes
<title>	Title	plain text
<enclosure url> (itorrents)	InfoHash	extract [A-F0-9]{40} from URL; build magnet
<size>	SizeBytes	integer bytes
	Seeders/Leechers	

RSS element	provider.SearchItem	Notes
Seeds: N , Leechers M in <description>		regex-parse the description string e.g. 12 Dec 2025
<pubDate>	Date	10:15:20 +0200 (non- standard, no weekday)

VERDICT: VIABLE (lowest of the three). Genuine filter, passive CF, hash + seeders recoverable — but two friction points lower its rank: (a) hard dependency on a .lol → .fun redirect on a rotating-TLD host (the most fragile of the three); (b) hash is indirect (URL-embedded, uppercase hex) and seeders/leechers need free-text regex out of <description>. Acceptable as a later addition, behind TorrentDownloads + TokyoTosho.

REJECTED candidates (each with captured evidence)

GloDLS (glodls.to → gtorrents.site) — REJECTED (EZTV-class no-op)

- GET https://glodls.to/rss.php?q=<query> → 302 → gtorrents.site/rss.php, final HTTP 200, application/xhtml+xml.
- **Filter is a NO-OP** (the canonical EZTV trap): q=ubuntu and q=debian returned the **byte-identical 50-item "latest torrents" feed** — <guid> lists diff'd IDENTICAL; 0/50 titles contained the query term in either feed (first item for both was "Iron Maiden - Virtual XI ... [FLAC]"). The q parameter is ignored. **Fails bar #1.**

Torrent-Paradise (.ml / .org / .eu) — REJECTED (dead / parked)

- torrent-paradise.ml/api/search?q=ubuntu → HTTP 200 but content-type: text/html serving a **Hungarian online-casino spam page** ("Magyar online kaszinók 2026 ...") — the domain has been repurposed; the JSON API is gone. torrent-paradise.org and torrent-paradise.eu → curl (6) Could not resolve host. **Fails bar #1/#5 (no API).**

AcademicTorrents (academictorrents.com) — REJECTED (no per-query search feed)

- The site explicitly **blocks browse-scraping** (browse.php → 307 → checkb.htm, an HTML page that says: "Hello if you are reading this you

may be trying to scrape the browse page ... we instead ask you to search an XML file (in RSS format)". The offered feeds — `rss.xml` (recent) and `database.xml` (full DB dump) — take **NO query parameter**: `rss.xml` returns a fixed 44 KB "Recent Torrents" list regardless of input. `apiv2/.../search` → 404. There is no anonymous free-text-filtering endpoint. **Fails bar #1.**

Anidex (anidex.info) — REJECTED (unstable / down)

- GET `https://anidex.info/rss/?q=naruto` → first attempt curl (28) timeout (15 s, 0 bytes); retry with a browser UA → HTTP 502 (Cloudflare error page "Error 502"). The origin is unreachable/unstable. **Fails bar #3/#5.**

ext.to (ext.to, ExtraTorrent mirror) — REJECTED (403)

- GET `https://ext.to/rss/?q=ubuntu` → HTTP 403 (text/html). Server-side GET is refused. **Fails bar #3.**

iDope (idope.se) — REJECTED (dead)

- GET `https://idope.se/torrent-list/ubuntu/?r=1` → curl (28) connection timeout (12 s). Host unreachable. **Fails bar #5.**

Snowfl (snowfl.com) — REJECTED (token-gated, not anonymous-stable)

- GET `https://snowfl.com/index1.php?ubuntu<token>&...` → HTTP 404 "Error". Snowfl's search requires a per-session token minted by its JS bundle; there is no stable anonymous query endpoint a plain Go GET can hit. **Fails bar #2/#5.**

Bitmagnet public instances — REJECTED (no stable public host)

- `bitmagnet.io/graphql` → 404; `demo.bitmagnet.io` → DNS does not resolve. Bitmagnet is self-host software, not a public aggregator with a stable anonymous endpoint to pin. **Fails bar #5.**

BTDigg (btdig.com) — REJECTED (unreachable)

- GET `https://btdig.com/search?q=ubuntu&api=1` → curl (28) timeout (12 s). **Fails bar #5.**
-

Ranked VIABLE shortlist (ready to implement)

Rank	Provider	Endpoint	info_hash	Seeders/ Leechers
1	TorrentDownloads	GET ... torrentdownloads.pro/ rss.xml? type=search&search=<q>	native <info_hash> 40-hex	YES (native <seeders>/ <leechers>)
2	Tokyo Toshokan	GET ...tokyotosho.info/ rss.php?terms=<q>	magnet (base32 btih) in <description>	NO (absent)
3	LimeTorrents	GET ...limetorrents.lol/ searchrss/<q>/ (→ .fun)	indirect (40- hex in itorrents enclosure URL)	YES (regex from <description>)

All three: anonymous, passive-Cloudflare (no cf-mitigated challenge), RSS, implementable against the internal/provider/curated/nyaa RSS pattern.

Single best next-to-implement

TorrentDownloads (torrentdownloads.pro). Strongest evidence line:

search=ubuntu and search=debian each return 50 RSS items, 50/50 titles on-topic, **0 overlap**, with a native <info_hash> 40-hex element plus <seeders>/<leechers>/<size>/<pubDate> — and server: cloudflare with cf-mitigated ABSENT, served identically to a plain LavaServerBot/1.0 UA (no browser/UA gating).

It is the only round-2 candidate that ships a complete native field set (hash + seeders + leechers + size + date) in the feed, making it the cleanest drop-in clone of the existing Nyaa provider.