

Additional Curated Provider Candidates — Round 2 Live-Probe Research (2026-06-13)

Round-2 research feeding **Defect B** (on-device api-app embeds curated Go providers in `lava-api-go/internal/provider/curated/`).

Already shipped (6) — NOT re-proposed here: The Pirate Bay (`apibay.org`), YTS (`yts.mx`), Torrents-CSV (`torrents-csv.com`), BitSearch (`bitsearch.eu`), Knaben (`api.knaben.org`), Nyaa (`nyaa.si` RSS).

Already rejected with evidence (round 1 + prior) — NOT re-probed: EZTV (keyword no-op), 1337x (CF challenge), TorrentGalaxy (DNS dead), SolidTorrents (= BitSearch backend), Torlock (HTML timeout), `torrent-api-py` (404).

All probes were live, server-side `curl` GETs from a macOS host on 2026-06-13 (~09:38–11:50 UTC), bounded ≤ 18 s each, default UA Mozilla/5.0 (compatible; LavaResearchBot/1.0) unless noted. Per §11.4.6 no-guessing vocabulary: only captured facts are stated; anything not directly probed is marked UNCONFIRMED:.

The bar (ALL must hold, each PROVEN with a captured probe)

1. Free-text query genuinely FILTERS — two distinct queries return DIFFERENT result sets that each contain the query term (the EZTV no-op test).
 2. Anonymous — no API key, no login, no cookie.
 3. NOT Cloudflare-challenge-gated for a plain server-side Go GET (no HTTP 403 + `cf-mitigated: challenge` / "Just a moment").
 4. Returns an `info_hash` or magnet.
 5. Stable JSON (preferred) or RSS / trivially parseable response.
-



Candidate results

1. TorrentDownloads (torrentdownloads.pro) — VIALE (TOP PICK)

- **Search endpoint:** GET `https://www.torrentdownloads.pro/rss.xml?type=search&search=<query>` (a `www.torrentdownloads.pro` host; the bare host issues a redirect — follow it / pin `www.`).
- **Live probe:** HTTP 200, body is RSS XML (`<?xml ... <rss version='2.0'>`), bytes≈26627 for `q=ubuntu`. server: cloudflare present but **cf-mitigated ABSENT** — passive CDN, NOT a challenge gate. Works with a plain non-browser UA (LavaServerBot/1.0 → 50 items), so no UA-gating.
- **Auth:** anonymous; no key/cookie.
- **info_hash / magnet:** YES — dedicated `<info_hash>` element, 40-hex lowercase (e.g. `19987310611f0fae3b2c9678601b92967d938aff`). Magnet is built client-side from the hash + the curated public-tracker list (same pattern as the Nyaa provider).

2-query filter-difference proof (the EZTV no-op test) — PASSES decisively:

Q1 `search=ubuntu` → 50 items; 50/50 `<title>` contain "ubuntu"
("Ubuntu 12 10 Server 64 Bit", "Ubuntu 10 10 i386 v14 2 5", ...)
Q2 `search=debian` → 50 items; 50/50 `<title>` contain "debian"
overlap between the two title sets: 0

Field mapping (RSS `<item>` → `provider.SearchItem`):

RSS element	<code>provider.SearchItem</code>	Notes
<code><title></code>	Title	plain text
<code><info_hash></code>	InfoHash	40-hex lowercase; build magnet from it
<code><size></code>	SizeBytes	integer bytes (e.g. 24898149)
<code><seeders></code>	Seeders	integer (e.g. 978)
<code><leechers></code>	Leechers	integer (e.g. 30)
<code><pubDate></code>	Date	RFC-1123 (Mon, 21 Oct 2013 13:20:49 +0200)
<code><link></code>	(detail URL)	site-relative <code>/torrent/<id>/<slug></code>

VERDICT: VIALE. Best round-2 candidate — full field set (the ONLY round-2 candidate with native `<info_hash>` + `<seeders>` + `<leechers>` + `<size>` all in the feed), clean filter, passive CF, no UA-

gating. General-purpose corpus (complements the niche providers). Maps cleanly onto the nyaa RSS provider shape.
Caveat (UNCONFIRMED): .pro is one of several rotating TorrentDownloads mirror TLDs; long-term host stability of .pro specifically is not proven beyond today's probe.

2. Tokyo Toshokan (tokyotosho.info) — VIABLE (anime/niche, no seeders)

- **Search endpoint:** GET `https://www.tokyotosho.info/rss.php?terms=<query>`
- **Live probe:** HTTP 200, content-type: application/rss+xml; charset=utf-8. server: cloudflare, cf-mitigated ABSENT — passive CDN.
- **Auth:** anonymous.
- **info_hash / magnet:** YES — each item's <description> (CDATA) embeds a full magnet:?xt=urn:btih:<BASE32-40>&tr=... link (e.g. magnet:?xt=urn:btih:AIHBWH5HWF2UQCHMDT2D306TP63HBLNX&tr=...). Note the btih is **base32** (32 chars), not 40-hex — the provider must accept the magnet as-is OR base32-decode → hex.

2-query filter-difference proof — PASSES:

Q1 terms=naruto → 139 items; first content titles on-topic
("Cutie Kuromi - Naruto chan.zip", "... (NARUTO -ナルト-).zip", ...)
Q2 terms=bleach → 146 items; on-topic
("Bleach The new Gotei-13 law - Final [ENG].zip", "... (Bleach)...", ...)
disjoint, query-specific result sets.

Field mapping (RSS <item> → provider.SearchItem):

RSS element	provider.SearchItem	Notes
<title>	Title	plain text
magnet: URL inside <description>	InfoHash / magnet	base32 btih; regex-extract from CDATA
Size: NNN MB inside <description>	SizeBytes	human-readable; must parse "113.89MB"
<pubDate>	Date	RFC-1123 GMT
<category>	(category)	e.g. "Other", "Hentai (Manga)"
seeders / leechers	— (ABSENT)	NOT in the feed → emit 0/unknown

VERDICT: VIABLE. Honest CapSearch (genuine filter), anonymous, magnet present. **Two caveats:** (a) NO seeders/leechers in the feed (must surface as 0/unknown — UI honesty); (b) btiH is base32, and size needs human-string parsing. Niche (anime/Japanese media) — overlaps Nyaa's domain but draws a distinct corpus + different uploaders. Lower priority than TorrentDownloads because of the missing seeders + base32 + the adult-heavy default result mix.

3. LimeTorrents (limetorrents.lol → .fun) — VIABLE (redirect + indirect hash)

- **Search endpoint:** GET `https://www.limetorrents.lol/searchrss/<query>/` — issues HTTP 301 → `https://www.limetorrents.fun/searchrss/<query>/` (final host serves the RSS). **Must follow the redirect** (the .lol/.fun TLD rotates; pinning a single TLD is fragile).
- **Live probe (final host):** HTTP 200, RSS XML, server: cloudflare, cf-mitigated ABSENT — passive CDN.
- **Auth:** anonymous.
- **info_hash / magnet:** INDIRECT — no `<info_hash>` element and no magnet; the 40-hex hash is embedded in the `<enclosure url>` pointing to `https://itorrents.net/torrent/<40HEX>.torrent` (e.g. `itorrents.net/torrent/232CD67EB3FFBD7C37BF9EC3EE887417E5AE1EE6.torrent`). Must regex-extract the `[A-F0-9]{40}` and build the magnet client-side.

2-query filter-difference proof — PASSES:

Q1 `searchrss/ubuntu/` → 40 items; 41 `<title>` lines incl. ubuntu (41 = 40 items + feed header)

Q2 `searchrss/debian/` → 40 items; 41 `<title>` lines incl. debian
overlap between the two title sets: 0

Field mapping (RSS `<item>` → `provider.SearchItem`):

RSS element	provider.SearchItem	Notes
<code><title></code>	Title	plain text
<code><enclosure url></code> (itorrents)	InfoHash	extract <code>[A-F0-9]{40}</code> from URL; build magnet
<code><size></code>	SizeBytes	integer bytes
Seeds: N , Leechers M in <code><description></code>	Seeders/Leechers	regex-parse the description string
<code><pubDate></code>	Date	e.g. 12 Dec 2025 10:15:20 +0200 (non-

RSS element	provider.SearchItem	Notes
		standard, no weekday)

VERDICT: VIABLE (lowest of the three). Genuine filter, passive CF, hash + seeders recoverable — but two friction points lower its rank: (a) hard dependency on a .lol → .fun redirect on a rotating-TLD host (the most fragile of the three); (b) hash is indirect (URL-embedded, uppercase hex) and seeders/leechers need free-text regex out of <description>. Acceptable as a later addition, behind TorrentDownloads + TokyoTosho.

REJECTED candidates (each with captured evidence)

GloDLS (glodls.to → gtorrents.site) — REJECTED (EZTV-class no-op)

- GET https://glodls.to/rss.php?q=<query> → 302 → gtorrents.site/rss.php, final HTTP 200, application/xhtml+xml.
- **Filter is a NO-OP** (the canonical EZTV trap): q=ubuntu and q=debian returned the **byte-identical 50-item "latest torrents" feed** — <guid> lists diff'd IDENTICAL; 0/50 titles contained the query term in either feed (first item for both was "Iron Maiden - Virtual XI ... [FLAC]"). The q parameter is ignored. **Fails bar #1.**

Torrent-Paradise (.ml / .org / .eu) — REJECTED (dead / parked)

- torrent-paradise.ml/api/search?q=ubuntu → HTTP 200 but content-type: text/html serving a **Hungarian online-casino spam page** ("Magyar online kaszinók 2026 ...") — the domain has been repurposed; the JSON API is gone. torrent-paradise.org and torrent-paradise.eu → curl (6) Could not resolve host. **Fails bar #1/#5 (no API).**

AcademicTorrents (academictorrents.com) — REJECTED (no per-query search feed)

- The site explicitly **blocks browse-scraping** (browse.php → 307 → checkb.htm, an HTML page that says: *"Hello if you are reading this you may be trying to scrape the browse page ... we instead ask you to search an XML file (in RSS format)"*). The offered feeds — rss.xml (recent) and database.xml (full DB dump) — take **NO query parameter**: rss.xml returns a fixed

44 KB "Recent Torrents" list regardless of input. `apiv2/.../search` → 404. There is no anonymous free-text-filtering endpoint. **Fails bar #1.**

Anidex (anidex.info) — REJECTED (unstable / down)

- GET `https://anidex.info/rss/?q=naruto` → first attempt curl (28) timeout (15 s, 0 bytes); retry with a browser UA → HTTP 502 (Cloudflare error page "Error 502"). The origin is unreachable/unstable. **Fails bar #3/#5.**

ext.to (ext.to, ExtraTorrent mirror) — REJECTED (403)

- GET `https://ext.to/rss/?q=ubuntu` → HTTP 403 (text/html). Server-side GET is refused. **Fails bar #3.**

iDope (idope.se) — REJECTED (dead)

- GET `https://idope.se/torrent-list/ubuntu/?r=1` → curl (28) connection timeout (12 s). Host unreachable. **Fails bar #5.**

Snowfl (snowfl.com) — REJECTED (token-gated, not anonymous-stable)

- GET `https://snowfl.com/index1.php?ubuntu<token>&...` → HTTP 404 "Error". Snowfl's search requires a per-session token minted by its JS bundle; there is no stable anonymous query endpoint a plain Go GET can hit. **Fails bar #2/#5.**

Bitmagnet public instances — REJECTED (no stable public host)

- `bitmagnet.io/graphql` → 404; `demo.bitmagnet.io` → DNS does not resolve. Bitmagnet is self-host software, not a public aggregator with a stable anonymous endpoint to pin. **Fails bar #5.**

BTDigg (btdig.com) — REJECTED (unreachable)

- GET `https://btdig.com/search?q=ubuntu&api=1` → curl (28) timeout (12 s). **Fails bar #5.**
-

Ranked VIABLE shortlist (ready to implement)

Rank	Provider	Endpoint	info_hash	Seeders/Leechers
1	TorrentDownloads	GET ... torrentdownloads.pro/ rss.xml? type=search&search=<q>	native <info_hash> 40-hex	YES (native <seeders>/ <leechers>)
2	Tokyo Toshokan	GET ...tokyotosho.info/ rss.php?terms=<q>	magnet (base32 btih) in <description>	NO (absent)
3	LimeTorrents	GET ...limetorrents.lol/ searchrss/<q>/ (→ .fun)	indirect (40- hex in itorrents enclosure URL)	YES (regex from <description>)

All three: anonymous, passive-Cloudflare (no cf-mitigated challenge), RSS, implementable against the internal/provider/curated/nyaa RSS pattern.

Single best next-to-implement

TorrentDownloads (torrentdownloads.pro). Strongest evidence line:

search=ubuntu and search=debian each return 50 RSS items, 50/50 titles on-topic, **0 overlap**, with a native <info_hash> 40-hex element plus <seeders>/<leechers>/<size>/<pubDate> — and server: cloudflare with cf-mitigated ABSENT, served identically to a plain LavaServerBot/1.0 UA (no browser/UA gating).

It is the only round-2 candidate that ships a complete native field set (hash + seeders + leechers + size + date) in the feed, making it the cleanest drop-in clone of the existing Nyaa provider.