

scripts/run-helixqa-challenges.sh — User Guide

Last verified: 2026-05-16 (Phase 4 follow-up B — 4 open questions resolved) **Inheritance:** HelixConstitution §11.4.18 (script documentation mandate) + Lava §6.AE (Comprehensive Challenge Coverage) + §6.J (Anti-Bluff Functional Reality) + §6.X (Container-Submodule Emulator Wiring — implemented via --runner=containerized) + §6.W (mirror-host boundary — per-script audit at docs/helixqa-script-audit.md) **Classification:** project-specific (the per-script wiring choices + the §6.W audit results are Lava-specific; the host-vs-container delegation is universal per §6.X)

Overview

Thin Lava-side wrapper around HelixQA's 11 Challenge scripts shipped at submodules/helixqa/challenges/scripts/. Invokes them in sequence, captures per-script stdout/stderr to log files, writes a roll-up attestation JSON, and reports per-script PASS / FAIL / SKIP outcomes.

Implements **Option 1** from docs/plans/2026-05-16-helixqa-integration-design.md — shell-level wiring with NO modification to HelixQA. Future cycles MAY:

- Add Go-package linking for Group A packages (Option 2 in the design).
- Migrate Compose UI Challenge Tests to use HelixQA as backend (Option 3 in the design — deferred indefinitely).

The 11 wired scripts

Canonical list (matches HELIXQA_SCRIPTS array in the script body):

#	Script	What it checks (HelixQA's own framing)	Toolchain
1	anchor_manifest_challenge.sh		none

#	Script	What it checks (HelixQA's own framing)	Toolchain
		docs/ behavior-anchors.md rows + per-anchor file/ symbol resolution	
2	bluff_scanner_challenge.sh	scanner self-test + tree-wide bluff scan	go
3	chaos_failure_injection_challenge.sh	chaos-engineering harness	none
4	ddos_health_flood_challenge.sh	DDoS-resilience health-endpoint flooding	none
5	host_no_auto_suspend_challenge.sh	host suspend / hibernate / sleep targets masked + sleep.conf override (Linux-only)	none
6	mutation_ratchet_challenge.sh	go-mutesting kill-rate ratchet	go
7	no_suspend_calls_challenge.sh	source-tree scan for forbidden host-power-management invocations	none
8	scaling_horizontal_challenge.sh	horizontal-scaling load test	none
9	stress_sustained_load_challenge.sh	sustained-load stress test	none

#	Script	What it checks (HelixQA's own framing)	Toolchain
10	ui_terminal_interaction_challenge.sh	terminal UI interaction (limited Lava applicability)	none
11	ux_end_to_end_flow_challenge.sh	end-to-end UX flow harness	none

Toolchain column source: HELIXQA_TOOLCHAIN_MAP constant in the wrapper. go means the script invokes a Go binary (go-mutesting, the bluff scanner) that requires go on PATH.

Usage

```
# Run ALL 11 HelixQA Challenge scripts on the HOST (workstation iteration)
bash scripts/run-helixqa-challenges.sh

# Run only specific scripts (comma-separated basenames without .sh)
bash scripts/run-helixqa-challenges.sh --only host_no_auto_suspend_challenge,no_suspend_calls_challenge

# Custom evidence dir
bash scripts/run-helixqa-challenges.sh --evidence-dir .lava-ci-evidence/myrun/helixqa

# Custom evidence dir via env var (useful when wrapping)
LAVA_HELIXQA_EVIDENCE_DIR=/tmp/helixqa bash scripts/run-helixqa-challenges.sh

# Halt the loop on first non-zero exit (default: continue)
bash scripts/run-helixqa-challenges.sh --stop-on-fail

# Suppress stdout summary (machine consumption only)
bash scripts/run-helixqa-challenges.sh --json-only

# $6.X containerized runner (gate-mode evidence)
bash scripts/run-helixqa-challenges.sh --runner containerized --container-image docker.io/library/golang:1.22

# Skip Go-requiring scripts when Go is unavailable
bash scripts/run-helixqa-challenges.sh --require-toolchain none
```

Inputs

Arg / Env	Description
--evidence-dir <dir>	Output directory (default: \$LAVA_HELIXQA_EVIDENCE_DIR if set; else .lava-ci-evidence/helixqa-challenges/<UTC-timestamp>/)
LAVA_HELIXQA_EVIDENCE_DIR env	Env-var alternative to --evidence-dir. CLI flag wins if both are set.
--only NAME1,NAME2,...	Run only the listed scripts (basename without .sh); errors if none match
--continue-on-fail	Keep running scripts after a FAIL (default behavior)
--stop-on-fail	Halt the loop on first FAIL (useful for triage of cascading failures)
--json-only	Suppress the human-readable stdout summary; only write the attestation JSON
--runner host containerized	\$6.X delegation: host (default — workstation iteration) or containerized (required for \$6.AE gate runs). containerized honest-fail-fasts (exit 4) when submodules/containers is absent.
--container-image <image>	Container image for --runner=containerized. Default: docker.io/library/golang:1.22 (provides Go toolchain for HelixQA's Go-requiring scripts).
--container-runtime podman docker	Container runtime for --runner=containerized. Default: auto-detect (podman preferred per \$6.U).
--require-toolchain go none	\$6.J toolchain precondition. go (default): wrapper exits 4 if go is absent from PATH (host mode) or always available (container mode). none: Go-requiring scripts (bluff_scanner, mutation_ratchet) SKIP with a clear precondition message when go is absent.

Outputs

<evidence-dir>/
└─ helixqa-attestation.json # roll-up: pin, runner,

```

totals, per-script rows
├─ anchor_manifest_challenge.log      # per-script stdout+stderr
├─ bluff_scanner_challenge.log
├─ chaos_failure_injection_challenge.log
├─ ddos_health_flood_challenge.log
├─ host_no_auto_suspend_challenge.log
├─ mutation_ratchet_challenge.log
├─ no_suspend_calls_challenge.log
├─ scaling_horizontal_challenge.log
├─ stress_sustained_load_challenge.log
├─ ui_terminal_interaction_challenge.log
└─ ux_end_to_end_flow_challenge.log

```

Attestation JSON shape

```

{
  "timestamp": "2026-05-16T03:52:15Z",
  "helixqa_pin": "403603dbd47b971549709bd9f2efa30dca810097",
  "helixqa_runner": "host",
  "helixqa_runner_caveat": "$6.X container-wrapping NOT in use
    this run; rerun with --runner=containerized for $6.AE
    gate-mode evidence",
  "require_toolchain": "go",
  "toolchain_available": true,
  "container_runtime": "n/a",
  "container_image": "n/a",
  "w_excluded_count": 0,
  "w_exclusion_audit": "docs/helixqa-script-audit.md",
  "total_scripts": 11,
  "pass_count": 5,
  "fail_count": 2,
  "skip_count": 4,
  "halted_early": false,
  "all_passed": false,
  "scripts": [
    {
      "name": "anchor_manifest_challenge.sh",
      "result": "FAIL",
      "exit_code": 1,
      "duration_seconds": 0,
      "log": "anchor_manifest_challenge.log"
    }
  ]
}

```

When `--runner=containerized` is selected the `helixqa_runner` field reads `containerized`, the caveat is replaced with `"(no caveat – $6.X containerized runner active)"`, and `container_runtime` / `container_image` carry the actual values used.

Exit codes

Code	Meaning
0	

Code	Meaning
	All invoked scripts returned 0 (or SKIP exit 2). Zero FAIL outcomes.
1	One or more invoked scripts returned non-zero exit other than the SKIP-canonical 2. Real defect surfaced by HelixQA OR --stop-on-fail halted the loop.
2	Invalid arguments (--only matched no scripts; unknown flag; invalid --runner or --require-toolchain value).
3	Missing dependency — submodules/helixqa directory absent, OR the canonical 11-script wired list drifted from what the pin actually ships.
4	Missing runtime — --runner=containerized without submodules/containers bootstrapped, OR --runner=containerized without a container runtime on PATH, OR --require-toolchain=go without go on PATH.

The 0/1/2/3/4 split is **anti-bluff load-bearing**: exit 0 means scripts genuinely ran with no failures; exit 3 means we honestly cannot claim to have run them (precondition gap); exit 4 means a runtime / toolchain dependency is missing (silent degradation forbidden per §6.J). Per §6.J we never silently mask a precondition gap as "success".

Per-script SKIP / FAIL semantics

The wrapper treats each script's own exit codes as authoritative and classifies:

- exit 0 → **PASS** — script reports success
- exit 2 → **SKIP** — script's documented "scanner missing / precondition unmet" exit code (per `no_suspend_calls_challenge.sh` convention)
- any other non-zero → **FAIL** — real defect OR Linux-only check failing on macOS/other host

Additionally the wrapper itself may classify a script as SKIP BEFORE invoking it:

- The script is on the `HELIXQA_W_EXCLUSIONS` array (§6.W audit found a violation) → **SKIP** exit-2 with a per-script log line citing the audit doc
- The script declares `toolchain: go` in `HELIXQA_TOOLCHAIN_MAP` AND `--require-toolchain=none` AND `go` absent from `PATH` → **SKIP** exit-2 with a precondition-unmet message

Many of the 11 HelixQA scripts are LINUX-SPECIFIC (they invoke `systemctl`, `read /etc/systemd/`, `parse journalctl`). On macOS or any host without `systemd`, these will FAIL — not because Lava is broken, but because the HelixQA author wrote them assuming a

Linux host. **This is the honest behavior** — per §6.J the wrapper does NOT mis-classify "wrong-OS" as PASS or as SKIP without operator intervention. Operators on macOS hosts SHOULD use `--only` to filter to portable scripts (anchor_manifest_challenge, bluff_scanner_challenge, mutation_ratchet_challenge — when the HelixQA-internal scanner directory is present), OR use `--runner=containerized` to provide a Linux execution surface inside a podman/docker container.

§6.X containerized runner

When invoked with `--runner=containerized`, the wrapper:

1. Verifies submodules/containers is bootstrapped (else exit 4)
2. Auto-detects the container runtime (podman preferred; falls back to docker)
3. Selects the container image (default: docker.io/library/golang:1.22 to provide the Go toolchain HelixQA's Go-requiring scripts need)
4. For each script, invokes `<runtime> run --rm --user $(id -u):$(id -g) -v $REPO_ROOT:$REPO_ROOT:rw -w $HELIXQA_SCRIPTS_DIR -e HOME=/tmp $CONTAINER_IMAGE bash ./<script>`
5. Writes the same per-script logs + attestation JSON to the evidence dir on the host filesystem (via the bind mount)

The container runs as the host UID so evidence files retain operator ownership — no root-owned outputs.

Limitations of the containerized runner (honestly disclosed):

- Scripts that probe HOST systemd state (host_no_auto_suspend_challenge.sh) will see the CONTAINER's systemd state, not the host's. Their results inside a container are NOT meaningful for the host the operator runs on; treat them as HelixQA's own self-test rather than as a Lava-host audit.
- Scripts that connect to LOCAL services (chaos, ddos, scaling, stress against localhost:8081) will hit the container's localhost, not the host's. The operator MUST either run the local service inside the same network namespace OR override HELIXQA_HEALTH_URL to a reachable address. The wrapper does NOT pre-validate this — it is operator-supplied per §6.W audit.

§6.W per-script mirror-policy audit

docs/helixqa-script-audit.md is the source-of-truth audit. As of the last audit (2026-05-16), 0 of 11 scripts violate §6.W on default config — the HELIXQA_W_EXCLUSIONS array is therefore empty. The audit MUST be re-run when the submodules/helixqa pin bumps; the

wrapper's wiring-drift check (exit 3 on added/removed scripts) prompts re-audit but does NOT mechanize the per-script grep — that remains a human-driven audit per §6.J ("real grep results only — no manufactured findings").

Wiring into scripts/run-challenge-matrix.sh

The §6.AE Challenge matrix runner accepts an opt-in `--include-helixqa` flag that invokes this wrapper as a sibling step BEFORE the AVD matrix runs:

```
bash scripts/run-challenge-matrix.sh --include-helixqa
```

The flag is **OFF by default** so existing matrix invocations (operator iterations, prior CI flows) are unaffected. When enabled:

- HelixQA wrapper runs on the host FIRST (host independent of AVD matrix)
- Wrapper evidence lands at `<matrix-evidence-dir>/helixqa/`
- A non-zero HelixQA exit is folded into the matrix runner's final exit code (matrix exit dominates if both fail; HelixQA exit promotes matrix-0 to 1 if HelixQA reports FAIL)
- HelixQA does NOT run inside the §6.X-debt darwin/arm64 gate-host skip — it runs regardless because it doesn't need KVM

To force the matrix-aggregated HelixQA invocation to use the containerized runner, the operator currently runs the wrapper separately:

```
bash scripts/run-helixqa-challenges.sh --runner containerized
bash scripts/run-challenge-matrix.sh           # without --include-helixqa
```

A future cycle MAY teach the matrix runner to pass `--runner=containerized` through; today the matrix runner invokes the wrapper with default arguments (host mode) since the matrix runner itself only fully runs on Linux x86_64 hosts where the host runner is effectively equivalent.

Anti-bluff posture

Per docs/plans/2026-05-16-helixqa-integration-design.md §"6.J anti-bluff posture":

1. **Each invoked script MUST actually be invocable.** The wrapper's pre-flight rejects (exit 3) if the wired list drifts from the actual pin contents.
2. **Each result MUST be from a real execution.** Per-script log files are written from the actual subprocess output — captured

stdout+stderr go to <evidence-dir>/<script>.log. The hermetic test tests/check-constitution/test_helixqa_wiring.sh asserts the fixture's stdout string appears in the produced log.

3. **Per §6.AC non-fatal telemetry:** SKIP exits do NOT cause the wrapper itself to FAIL but ARE recorded distinctly from PASS. The attestation JSON's skip_count is operator-visible so a regression that flips many PASSes to SKIPS is detectable.
4. **Per §6.J/§6.L:** the wrapper exists; HelixQA-emitted FAILs are real defects to triage; HelixQA-emitted SKIPS are honest acknowledgments of precondition gaps — NOT false-pass coverage.
5. **Per §6.X:** --runner=containerized honest-fail-fasts (exit 4) when submodules/containers is absent. The wrapper NEVER silently degrades containerized → host because that would be the exact bluff §6.X exists to prevent.
6. **Per §6.J:** --require-toolchain=go (default) exits 4 when go is absent rather than silently SKIPPING Go-requiring scripts. Silent-skip-with-default would produce false-green coverage; operators who knowingly work without Go MUST opt in via --require-toolchain=none.
7. **Per §6.W:** the per-script audit at docs/helixqa-script-audit.md is the binding source-of-truth for the HELIXQA_W_EXCLUSIONS array. Manufactured findings are forbidden; only real grep results may populate the audit.

Hermetic test

tests/check-constitution/test_helixqa_wiring.sh — 11 fixtures:

Fixture	Asserts
test_passes_when_helixqa_present_and_all_green	Wrapper exit 0; attestation reports 11/0/0; per-script log files exist + contain fixture stdout (providing subprocess actually ran)
test_fails_when_script_missing	Wrapper exit 3 (wiring drift); error message names the missing script
test_skip_mode_when_helixqa_absent	Wrapper exit 3 (missing submodule); error message includes the git submodule update --init submodules/helixqa remediation command; no attestation written
test_fail_classified_correctly	Wrapper exit 1 when script returns 1; attestation marks the

Fixture	Asserts
	script as FAIL (not S — anti-bluff classification) Q1: Wrapper exit 4; error message name
test_containerized_runner_fails_fast_when_containers_absent	\$6.X + git submodule update --init submod containers immediat
test_host_runner_default_attestation_shape	Q1: default-runner attestation has helixqa_runner: host caveat field + require_toolchain: g w_exclusion_audit po
test_require_toolchain_go_fails_fast_when_go_absent	Q2: Wrapper exit 4 + go absent + default require-toolchain=go error names the -- require-toolchain=no escape hatch
test_require_toolchain_none_skips_go_scripts_when_go_absent	Q2: Wrapper exit 0; requiring scripts (bluff_scanner, mutation_ratchet) marked SKIP; non-G scripts still PASS
test_evidence_dir_env_var_override	Q3: LAVA_HELIXQA_EVIDENC env var redirects attestation; default dated dir is NOT populated
test_invalid_runner_rejected	Q1: Wrapper exit 2 explicit error on unknown --runner va
test_invalid_require_toolchain_rejected	Q2: Wrapper exit 2 explicit error on unknown --require- toolchain value

Run: bash tests/check-constitution/test_helixqa_wiring.sh

Open questions (resolved 2026-05-16 — Phase 4 follow-up B)

The four open questions from the initial Option 1 commit were resolved in this commit:

1. **Q1 — Container vs host runner** (§6.X integration):
RESOLVED. New `--runner=host|containerized` flag (default: `host` for workstation iteration; `containerized` required for §6.AE gate runs). `Containerized` mode `honest-fail-fasts` (exit 4) when `submodules/containers` is absent — no silent degradation per §6.J.
2. **Q2 — Real-deps vs stub gating**: RESOLVED. New `--require-toolchain=go|none` flag (default: `go`). Per-script `toolchain` requirement declared in `HELIXQA_TOOLCHAIN_MAP` constant (`bluff_scanner` + `mutation_ratchet` need `Go`; the other 9 do not). `--require-toolchain=go` exits 4 when `go` absent; `--require-toolchain=none` SKIPS `Go`-requiring scripts with a clear precondition message.
3. **Q3 — Evidence-directory boundary**: RESOLVED. `--evidence-dir` flag is preferred; `LAVA_HELIXQA_EVIDENCE_DIR` env-var override added for parent-runner wrapping use cases. Default location remains `.lava-ci-evidence/helixqa-challenges/<UTC-timestamp>/` per Lava convention.
4. **Q4 — §6.W mirror policy**: RESOLVED. Per-script audit landed at `docs/helixqa-script-audit.md`; `HELIXQA_W_EXCLUSIONS` array constant in the wrapper consumes the audit's findings. 0/11 scripts currently violate §6.W on default config (the array is empty); re-audit owed on every `submodules/helixqa` pin bump.

See this commit's body for the per-question implementation details + falsifiability rehearsals.

Cross-references

- `docs/plans/2026-05-16-helixqa-integration-design.md` — design doc (Option 1 chosen this cycle; Option 2 + Option 3 deferred)
- `docs/scripts/run-challenge-matrix.sh.md` — the `--include-helixqa` flag
- `docs/helix-constitution-gates.md` — CM-HELIXQA-WIRING + CM-HELIXQA-§6.W-AUDIT gate rows
- `docs/helixqa-script-audit.md` — §6.W per-script audit (Phase 4 follow-up B Q4 resolution)
- `submodules/helixqa/CLAUDE.md` — HelixQA's own anti-bluff covenant
- `submodules/helixqa/README.md` — HelixQA framework overview
- `Lava CLAUDE.md` §6.AE + §6.J + §6.X + §6.W + §6.AD (HelixConstitution inheritance)