Issues_Summary

Open workable items (current_location = Issues), regenerated from the SQLite single-source-of-truth (§11.4.12).

Counts by Type × Status

Type	Status	Count
Bug	In progress	1
Bug	Operator-blocked	2
Bug	Queued	20
Bug	Ready for testing	1
Task	Operator-blocked	2
Task	Queued	7
TOTAL		**33**

Items

ATM ID	Type	Status	Severity	Description
LVA-008	Bug	Ready for testing	P1	ROOT CAUSE CORRECTED 2026-08-26 – supersedes the androidx-navigation attrib retained below. C06 + C11 (Challenge11ArchiveOrgAnonymousSearc crashed the app PROCESS at activity-des with IllegalStateException "State must be least CREATED to be moved to DESTROY the inner search/search_input entry. The c was NOT an upstream AndroidX bug and existed in AndroidX code. It was a Lava threading defect: under createAndroidComposeRule, FrameDeferringContinuationInterceptor resumed the Orbit collectSideEffect continuation OFF the main thread, so navHostController.navigate(...) executed main; navigation-runtime 2.9.1 inserts the back-stack entry BEFORE promoting its lifecycle, and LifecycleRegistry.setCurren

ATM ID	Type	Status	Severity	Description
				throws when invoked off-main — leaving t entry permanently INITIALIZED, which produced the illegal INITIALIZED-to-DESTROYED transition at Activity destroy FIXED 2026-06-30 by commit ccdd84c1, w wraps the navigate call in runOnMainThr core/navigation/src/main/kotlin/lava/navig NavigationController.kt:101. EVIDENCE CHAIN: last reproduction 2026-06-29 21:3 runOnMainThread guard added 2026-06-3 11:24; the same journey ran clean 2026-0 17:48; zero occurrences of the crash sign in any committed evidence since. DISPRO HYPOTHESIS (retained as forensic record 11.4.6, not erased): the item was attribute nested-NavHost lifecycle defect in androi navigation (b/244910446 family) and held Operator-blocked pending an upstream re Eight app-level candidate fixes were implemented and device-falsified: inner-NavHost Activity-scoped LifecycleOwner, to-outer-NavHost, ON_STOP force-pop (NavTeardownGuard), nav-compose 2.9.1 to-2.9.8-to-2.10.0-alpha bump, LenientTeardownRule, atomic popUpTo re launchSingleTop dedupe, and a full neste NavHost-architecture collapse. Each reproduced the identical ISE, and that uniformity was read as proof of an upstre defect. WHY THAT READING WAS WRON eight attempts changed WHERE navigate called or HOW the back stack was structu not one changed WHICH THREAD naviga ran on. The off-main call — the actual cau survived all eight, so all eight reproduced same crash. Uniform reproduction across fixes was evidence of an untouched comm cause, not of an upstream defect. The filin ready upstream repro package at docs/lva upstream-repro/ was DELETED 2026-08-2 filing it would have reported to AndroidX that does not exist in their code. RELATE CLEANUP: the autonomous-QA PASS-ove that converted this crash signature into a reported PASS (scripts/autonomous-qa/ru iteration.sh) was removed 2026-08-26 — v the cause fixed, it was a live 6.J bluff mec reporting a genuine crash as green. CLOS PENDING — THIS ITEM IS NOT CLOSED

ATM ID	Type	Status	Severity	Description
				closure requires a C06 + C11 device re-ru HEAD on a 6.AH-conformant container or emulator (never host-direct, never a live A device), both GREEN, with the JUnit XML captured as closure evidence.
LVA-121	Task	Queued	P3	UNCONFIRMED: this is a HYPOTHESIS, demonstrated defect. It was NOT reprodu during feature 002 and no failing case wa constructed. Filed so it is not lost; it must either reproduced (then re-filed as a Bug RED evidence) or explicitly ruled out. CLA scripts/pipeline/phase-01-build.sh counts ARTIFACT lines rather than distinct labels five duplicate labels would satisfy the >= PASS condition. WHY IT WAS NOT PROM TO A FINDING: not reachable through th sub-scripts, each of which prints a given l most once. To close, either construct a rea producing path or record the argument th none exists. Source of truth: specs/002-bu test-distribute-pipeline/progress.yml and 002-build-test-distribute-pipeline/tasks.m
LVA-122	Task	Queued	P3	UNCONFIRMED: this is a HYPOTHESIS, demonstrated defect. It was NOT reprodu during feature 002 and no failing case wa constructed. Filed so it is not lost; it must either reproduced (then re-filed as a Bug RED evidence) or explicitly ruled out. CLA the find ... 2>/dev/null inside recompute_evidence_summary in scripts/ pipeline/lib/run-report.sh would report ze Evidence Records for an unreadable phas directory, which finalize_run_report would treat as zero rejections. WHY IT WAS NO PROMOTED TO A FINDING: no realistic producing scenario was constructed. To c either construct one or record why the ph directory cannot become unreadable mid- Source of truth: specs/002-build-test-distr pipeline/progress.yml and specs/002-build distribute-pipeline/tasks.md.
LVA-123	Task	Queued	P3	UNCONFIRMED: this is a HYPOTHESIS, demonstrated defect. It was NOT reprodu during feature 002 and no failing case wa constructed. Filed so it is not lost; it must either reproduced (then re-filed as a Bug RED evidence) or explicitly ruled out. CLA the compose-file port parse in scripts/pipe

ATM ID	Type	Status	Severity	Description
				phase-04-live-verify-api.sh is a no-op whe LAVA_API_LISTEN value in docker-compo is unquoted. WHY IT WAS NOT PROMOT A FINDING: the real docker-compose.yml quotes the value, and a wrong port fails c at curl rather than passing vacuously. To either construct a producing case or mak parse fail loudly on an unrecognized shap Source of truth: specs/002-build-test-distr pipeline/progress.yml and specs/002-build distribute-pipeline/tasks.md.
LVA-128	Task	Operator-blocked	P0	WHAT: feature 002 task T054 is a MANDA code review of scripts/advance-all-submodules.sh before any invocation agai real submodule upstream, per plan.md Hu Checkpoint #2. This is the highest-blast-r new component in the feature: 16-plus submodule upstreams with automatic pin advancement, plus commit-and-push of ea submodule own local modifications. CURR STATE, stated honestly: the T050-T053 TI cycle against DISPOSABLE FIXTURE repo COMPLETE - 36/36 assertions green, independently re-run by the orchestrating session - and real submodules were confir untouched (git status --porcelain over submodules/ and constitution/ empty, git cached empty). scripts/advance-all-submodules.sh has NEVER been run agai real submodule or a real upstream. REVIE MUST COVER at minimum: the GOVERNANCE_DENY default-deny list ad by the wave-5 fix and the fact it is deliber not overridable via LAVA_ADVANCE_SUBMODULES; the 7-st ordering from research.md R-005 (fetch, compare, no-op or checkout-advance, reb and-test, discard-and-report on failure, co push the submodule own local modificatio update the parent pin); and the refusal pa REJECTED_BREAKING_CHANGE, REJECTED_PUSH_CONFLICT and REFUSED_GOVERNANCE_DENY. Source truth: specs/002-build-test-distribute-pipe progress.yml and specs/002-build-test-distribute-pipeline/tasks.md.
LVA-129	Task	Operator-blocked	P1	WHAT: feature 002 task T062 requires a f end-to-end run of quickstart.md Scenario phases, no shortcuts) on a DISPOSABLE

ATM ID	Type	Status	Severity	Description
LVA-131	Bug	Queued	P0	branch, per plan.md Human Checkpoint # before scripts/pipeline-build-test-distribut trusted to run against real master. CURRI STATE, stated honestly: the full end-to-en orchestrator run across all five wired pha has NOT been performed. It invokes Grad systemd and an Android emulator, and wa deliberately not run because it is this revi gate and because it would have contende parallel agents already using those toolch It is not claimed as done anywhere. SCOP NOTE: only phases precondition, build, te install_boot and live_verify are wired; dist docs and closure are deliberately not wire pending T040/T041 and T048/T049, and a for them via --until is a usage error rather a silent no-op. Source of truth: specs/002- test-distribute-pipeline/progress.yml and 002-build-test-distribute-pipeline/tasks.m WHAT: the phase-02 `go` test category caus lava-api-go's own §11.4.85 stress and cha suites to write their evidence into the repository-level .lava-ci-evidence/stress-cl jackett/ directory. Six files land there and are TRACKED (confirmed: `git ls-files` lis exactly 6 paths under that directory). The payload carries a captured_at stamp plus measured latency percentiles, so it differs every run by construction. Net effect: a SUCCESSFUL run left the working tree d and the next invocation was refused by th FR-000 precondition, which requires bran master AND a completely clean `git status porcelain` and exits 2 otherwise. EVIDENC measured directly: run 2026-08-23T10-17 finalized outcome PASS with phases precondition/build/test all PASS; the immediately following run 2026-08-23T10-31-21Z finalized outcome with a single precondition/FAIL phase, an phase-00 log reads verbatim `FR-000: precondition failed — working tree is not` This also falsifies quickstart.md Scenario claim that everything a run produces is gitignored. IMPACT: blocks FR-018 and S (a run must not dirty the working tree). STATUS: under active fix in the working t scripts/pipeline/phase-02-test-go.sh now e LAVA_STRESS_CHAOS_EVIDENCE_DIR

ATM ID	Type	Status	Severity	Description
				pointing at the run's own gitignored evide directory, and tests/pipeline/ test_phase_02_go_evidence_isolation.sh is present as regression coverage; neither is committed yet. CLOSURE CRITERIA: two consecutive full runs both reach their fina phase, with the tracked jackett files uncha after each.
LVA-132	Bug	Queued	P1	WHAT: scripts/pipeline-build-test-distribu installs `trap '_close_report' INT TERM` (line and then invokes each phase as a FOREGROUND pipeline: `"$SCRIPT_DIR/pip {script_name}" "$run_id" "$REPO_PATH_OVERR 2>&1 \| tee "$phase_log"` (lines 342 and 345 Bash does not execute a trap handler whi shell is blocked waiting on a foreground command; the handler is deferred until th command completes. A TERM delivered d a multi-minute Gradle build therefore doe finalize the report at signal time — it final only after the build finishes on its own, wh the opposite of what an interrupt is for. EVIDENCE, measured directly this sessio a minimal reproduction that mirrors the orchestrator's shape (trap on INT TERM, `sleep 6 | tee`): TERM was delivered at ep 1787681147 and the handler printed HAN RAN at epoch 1787681151 — a 4-second deferral exactly equal to the foreground pipeline's remaining duration. IMPACT: quickstart.md Scenario 5's assertion that TERM produces a finalized report does no while a phase is running, which is precise when an operator would send one. CLOSU CRITERIA: the phase runs in the backgro with the shell waiting on it interruptibly (equivalent seam), and a test delivers TER mid-phase and asserts the report is finaliz within a bounded interval rather than afte phase completes.
LVA-133	Bug	Queued	P2	WHAT: specs/002-build-test-distribute-pip quickstart.md Scenario 3 instructs the rea check out master, apply a deliberate one-l regression to production code, then invok `scripts/pipeline-build-test-distribute.sh` expect exit code 1 with a test-phase FAIL Evidence Record. That sequence cannot r the test phase: the mutation dirties the w tree, and scripts/pipeline/phase-00-

ATM ID	Type	Status	Severity	Description
LVA-137	Bug	Queued	P2	precondition.sh refuses any run whose `gi` `status --porcelain` is non-empty, printing `FR-000: precondition failed — working tre` `not clean` and exiting 2 (the branch and cl tree checks are at lines 39-48, both `exit 2` run therefore halts at precondition with e never reaches phase-02, and the scenario stated expected outcome is unreachable. the same quickstart already carries 2026- corrections to Scenario 3 on two OTHER (absent-versus-SKIPPED phases, and the vacuous zero-Distribution-Records asserti the document is maintained — this third c is simply not among the corrections. CLO CRITERIA: Scenario 3 is rewritten so its falsifiability demonstration is reachable — example by committing the mutation on a disposable branch, or by using the anti-bl Evidence Record variant the scenario's ov closing paragraph already describes as available now. WHAT: scripts/hooks/guard-forbidden-commands.sh line 208 gates the §6.T.3 for push check with `[[ "$COMMAND" =~ git([[:space:]]+[^[:space:]]+)*[[:space:]] +push ]]`. The expression is unanchored ar `([[:space:]]+[^[:space:]]+)*` group spans arbitrarily many arbitrary tokens, so it ma ANY command in which the literal git occ anywhere and the token push occurs later however unrelated the two are. Combined the standalone `-f` test on the next lines, a command that also contains a bare minus token (a routine file removal, for instance refused as a force-push. EVIDENCE, meas directly this session: the attempt to write ci-evidence/bluff-hunt/2026-08-23-cycle-g shaping-scripts.json was refused with the verbatim message `guardrails: BLOCKED — § force-push`, because the evidence file's ow prose describing this matcher contains th tokens. The write succeeded only by routi around the Bash tool entirely. The same matcher had already blocked the original 2026-08-23 attempt to record that hunt, w is why the evidence file was two days late ASSESSMENT: the guard fails SAFE — it refuses rather than permits — so no unsa operation resulted. The cost is different a

ATM ID	Type	Status	Severity	Description
				real: it trains agents to route around a saf guard to get ordinary work done, which e the guard's authority for the case it exists catch. CLOSURE CRITERIA: the matcher anchored to a genuine command position git invocation whose subcommand is push rather than to token co-occurrence anywh the string, with fixtures covering both a tr force-push and the benign co-occurrence above. RECURRENCE, 2026-08-26 (this it the pre-existing record; these are addition independent reproductions, filed here rat than as a duplicate item): the same match refused an ordinary documentation write TWICE more in a single session, on routin work that involved no version-control ope of any kind. In both instances the refusal triggered by the DOCUMENT'S OWN PRO the text under composition contained the git, a later occurrence of the token that n the rejected operation, and a bare minus- elsewhere in the same string. No comman invoked any version-control binary, and no in either write could have modified any re That brings the measured count of benign refusals by this matcher to at least three a two sessions, all of them on evidence-file documentation writes -- which is the popu most exposed to it, because documents AB the guard necessarily quote the very toke guard matches on. WHAT THIS CHANGES ABOUT THE ASSESSMENT: nothing abou direction of failure. The guard still fails SA and no unsafe operation has resulted. Wh recurrence establishes is that the false-po rate on documentation work is not inciden it is systematic and self-reinforcing, since attempt to WRITE UP the defect is itself b by the defect, which is why the original bl hunt evidence file landed two days late. T cost remains the one already recorded: ag learn to route around a safety guard to complete ordinary work, and a guard rout routed around has no authority left for the it exists to catch. The closure criteria abo unchanged; the fixture set should addition include a benign documentation string tha quotes all three trigger tokens and MUST permitted.

ATM ID	Type	Status	Severity	Description
LVA-148	Bug	Operator-blocked	P2	WHAT: the word channel names two diffe axes in this feature, and the conflict sits d under the repair path for LVA-120 and for section 6.AA clause 8 condition (D). MEASURED BY DIRECT INSPECTION: sp 002-build-test-distribute-pipeline/contract distribution-record.schema.json line 11 de channel as the Firebase APP, with an enu the client app distribution target and the app distribution target; scripts/firebase-distribute.sh lines 321 to 324 resolve a ch variable to release or debug, which is the BUILD VARIANT; and scripts/check-cycle-coverage.sh already receives the app-cha through its evidence-directory parameter, channel parameter is a SECOND axis wea the same word. That parameter is also IN and unvalidated: the variable is set by the argument-parsing branches at lines 37 an asserted non-empty at line 51, and never again in the remaining 128 lines of the sc both artifacts it checks resolve from the evidence directory and the version alone, debug, release and an arbitrary unknown all produce byte-identical output and exit RELATED AND SEPARATELY MEASURED release-channel evidence artifacts exist anywhere. The evidence tree under .lava-evidence/distribute-changelog partitions b only, into a client directory and an api-app directory, with no debug-versus-release sp any level. Implementing the channel para therefore requires first INVENTING what release-channel evidence IS, which is a de decision rather than a code fix. SCOPE NO this is not read as a duplicate: this item is vocabulary collision itself together with th absence of release-channel artifacts. The behavioural consequence inside firebase-distribute.sh is tracked at LVA-120, and cl condition (D) inherits both. DRAFTED OPTIONS: specs/002-build-test-distribute-pipeline/amendment-defects-proposal.md sections 2 and V-5.
LVA-149	Bug	In progress	P0	WHAT: scripts/pipeline-build-test-distribu resolves every git query and every eviden path against the CALLER'S working direc rather than against the repository the run supposed to be about. Three facts compos

ATM ID	Type	Status	Severity	Description
				the defect, each verified by direct inspect the script this session. First, the run repo commit identity is taken as commit_sha=" rev-parse HEAD)" with no -C "$REPO_RO so it names whatever repository the invok shell happens to be sitting in. Second, RU and _run_report_path() are RELATIVE pat under .lava-ci-evidence/pipeline-runs/, so report is deposited relative to the caller's directory, not the repository under test. T the script contains ZERO cd statements, s nothing ever reconciles the two: it never $REPO_ROOT, and it never re-resolves th evidence root against it. REPO_ROOT is computed and then not used for either pu MEASURED, not inferred: a run was invo against a fixture repository, and its FR-00 precondition guard passed legitimately, b the guard inspects the FIXTURE and the was on master with a clean tree. The run wrote its report into LAVA'S OWN .lava-ci evidence/pipeline-runs/ tree, carrying commit_sha 545920a160cc8e5591bb4c9c7a87793f37 -- Lava's HEAD, a commit the run never b never tested and never touched -- togethe outcome PASS and phases [(precondition, PASS)]. IMPACT: the constitution matches evidence to artifacts BY COMMIT SHA in separate places -- 6.Z's same-SHA test-evi rule, 6.AK's coverage-intersection gate, a 6.AA clause 8 condition (A), which refuses report whose commit_sha does not equal HEAD at distribute time. A report that na SHA it never exercised is therefore not m mislabelled: it is a false claim at rest, sitti the exact directory those three gates read the exact shape they accept. This is the 6. -- nothing was learned about that commit, recorded as though the commit passed. CONTAINMENT ALREADY PERFORMED such artifacts were preserved to scratchp forensics and removed from the evidence the 57 pre-existing legitimate reports wer untouched and were not modified in any w STATUS: a fix is IN FLIGHT with another in this session, which is why this item is fi progress rather than Queued -- it is record here so the defect cannot evaporate if tha

ATM ID	Type	Status	Severity	Description
				is interrupted (section 11.4.197). CLOSUI CRITERIA: every git invocation in the orchestrator is repository-scoped via -C "$REPO_ROOT" (or the script enters $REPO_ROOT before any of them), the ev root is resolved as an absolute path benea $REPO_ROOT, and a regression test invok orchestrator from a DIFFERENT working directory than the repository under test a asserts both that the report lands in the t repository's evidence tree and that its commit_sha equals the target repository's HEAD rather than the caller's.
LVA-151	Task	Queued	P1	WHAT: multiple agents working concurre ONE shared working tree cause constitut gates to report failures that are artifacts c contention rather than properties of the t FIVE separate gate failures during this se each dissolved on isolated re-measureme no code change was required for any of th MEASURED CASES, each re-measured in isolation before being dismissed: (a) a wo items check reported as a timeout was slo under concurrent load, not a hang -- it completed normally once the tree was qui a covenant-anchor check reported as a ra caused by the MEASURING process's own timeout terminating a probe part-way thr mutation, so the measurement destroyed state it was measuring; (c) a process-liven check reported a stale runner that was th pgrep pattern matching the checking proc itself, a self-match; (d) verify-all's markdo export-sync gate failed because a differen agent was regenerating .html and .pdf sib during the sweep, transiently leaving a .m newer than its own exports -- the isolated reported 282 files examined, 0 problems, and was repeated four consecutive times the identical result, which makes the true verify-all figure 58 of 58 rather than the 5 with-1-failure the contended sweep produ (e) check_constitution_test.sh's test_live_repo_passes failed twice under l and PASSED in isolation with exit 0 and 5 cases. COST: five separate diagnostic effo were spent establishing that nothing was wrong, and in case (d) a contended sweep produced a headline number that underst

ATM ID	Type	Status	Severity	Description
				the tree's real compliance -- a false RED, t erodes trust in the gates in exactly the wa false GREEN erodes trust in the tests. IM ON EVIDENCE: an attestation captured d contention records a verdict that is a prop of the scheduling, not of the tree, so it is r reproducible and cannot serve as evidenc either way. THE RULE THIS ITEM EXIST RECORD, stated plainly so a sixth diagnos not spent rediscovering it: A CONSTITUTIONAL-GATE FAILURE IS NO FINDING UNTIL IT REPRODUCES IN ISOLATION. A gate result captured while agents are writing to the tree is provision promoted to a finding only by a quiet-tree run, and it is that re-run's output that bel an evidence file. CLOSURE CRITERIA: th above is documented where gate-runners read it, and either a lock file serializes ga invocations across concurrent agents or e long-running gate acquires an advisory lo the duration of its sweep, with a test prov two concurrent invocations serialize rathe interleave.
LVA-152	Bug	Queued	P2	WHAT: tests/pipeline/ test_advance_all_submodules.sh defines a helper, run_script, whose body assigns the script's captured output to the GLOBAL v LAST_OUTPUT so that failure branches ca print what the script actually said. Every existing caller in cases 1 through 5 invoke helper inside a command substitution, in shape "$(run_script ...)". A command substitution executes in a SUBSHELL, an subshell's variable assignments do not propagate to the parent shell. LAST_OUT the calling shell therefore never receives value the helper assigned: it holds whatev held before -- an empty string on the first and thereafter the value left by whicheve earlier assignment last ran in the parent s CONSEQUENCE: every FAIL-branch diag in cases 1 through 5 that prints LAST_OU would print a stale value rather than the of the invocation that just failed. WHY IT NEVER BEEN VISIBLE: those FAIL branc not currently fire -- the suite passes -- so t stale value is never printed and nothing surfaces the defect. It is latent by constru

ATM ID	Type	Status	Severity	Description
				and will surface at exactly the worst mom when a genuine regression trips one of th branches and the operator reads a diagnc describing a DIFFERENT invocation than one that failed. SEVERITY RATIONALE, P assertion is wrong and no result is falsifie the suite's verdicts are correct. What is w the diagnostic attached to a future failure the cost is paid in misdirected debugging rather than in a wrong pass. PARTIAL STA ALREADY IN THE TREE: case 6, added d this session, does not use the helper's glo all -- it captures the script's output in the CURRENT shell instead, which is the corr pattern and can serve as the model. The c in cases 1 through 5 are unfixed. CLOSUI CRITERIA: run_script either returns its ou on stdout and every caller captures it in t current shell (the case 6 pattern), or the h is invoked without a command substitutio its global assignment reaches the caller; a one deliberately-failing fixture asserts tha printed diagnostic contains a string uniqu the failing invocation, which the current c cannot satisfy.
LVA-153	Task	Queued	P2	WHAT: the LVA-148 axis rename establish 'app' (Firebase app: client vs api-app) and 'build_variant' (release vs debug) as the tv distinct names. CHANGELOG_CHANNEL deliberately NOT renamed, because scrip pipeline/phase-05a-changelog-entry.sh:30 greps scripts/firebase-distribute.sh for the LITERAL string CHANGELOG_CHANNEL and two tests/pipeline/ files mirror that lit Renaming it in firebase-distribute.sh alon would break a live gate at a distance, whi why the rename agent stopped rather tha proceeding. OWED: one ATOMIC rename CHANGELOG_CHANNEL -> CHANGELO across 4 files / 5 sites, landed together. CONTAINED MEANWHILE: test case P5 the variable can only ever hold an app slu cannot silently acquire a build-variant val while the rename is pending. Source: self discovered during the LVA-148 rename, 2026-08-26.
LVA-154	Task	Queued	P2	WHAT: the LVA-148 rename separated 'ap from 'build_variant' in the contracts and s but the prose that describes them was no

ATM ID	Type	Status	Severity	Description
				updated in the same pass. Affected: CLAUDE.md 6.AA clause 8 still calls the h variant axis a channel; specs/002-build-te distribute-pipeline/spec.md likewise; and 002-build-test-distribute-pipeline/data-model.md's table now DESCRIBES A PROPERTY THAT NO LONGER EXISTS, s distribution-record.schema.json's channel renamed outright with no alias (0 records rest, so nothing became unreadable). Also at the same time: scripts/tag.sh:285-306 u channel_label / channel_dir for a THIRD s of the word, and scripts/pipeline/phase-02 hermetic.sh:72-75 cites tests/firebase/ repro_mode_both_channel_gap.sh, which deleted earlier the same day when combi distribute mode was retired. WHY IT MAT documentation describing a property that not exist sends the next reader to look for and the whole point of the rename was th word naming two axes is a trap. Governan edits require operator review, which is wh is filed rather than done. Source: self-discovered during the LVA-148 rename, 2026-08-26.
LVA-155	Bug	Queued	P1	WHAT: .githooks/pre-push Check 10 invok scripts/check-cycle-coverage.sh with --head=$sha where $sha IS the commit advancing the distribute pointer. The evid that commit references was written BEFO existed, so it cannot declare that commit's SHA. Now that the LVA-149/F2 fix makes commit-SHA binding actually bind (it was which is why 5 of 9 shipped cycles passed against an arbitrary wrong HEAD), this se reference surfaces: the next pointer-adva commit will be REJECTED at push time. MEASURED: the binding fix took wrong-H acceptance from 5/9 to 0/9; before the fix rejection was hidden. NOT closed by the f agent, deliberately and correctly: it declin implement an ancestor/code-identity esca because it could not prove the acceptance with a hermetic positive case, and an unte acceptance path inside a safety gate is pr the defect class it was dispatched to remo Three candidate options are written up in evidence log. DOES NOT BLOCK the curr push: Check 10 fires only on commits that

ATM ID	Type	Status	Severity	Description
				advance a distribute pointer, and the pend change set touches no last-version or dist changelog path (verified). Source: self- discovered during the LVA-149 fix, 2026-0 Needs an operator decision before the ne distribute.
LVA-156	Bug	Queued	P1	WHAT: re-running the REPAIRED cycle- coverage gate over all 9 shipped cycles sh hold up on substance — the gate was brok but those releases were not shipped on no 1.3.12-1077 is the exception and it FAILS the fixed gate: 2 of 3 claims uncovered, ar of the 13 covering refs its cycle-coverage- asserts have NO verdict record in the evi at all. One asserted ref is the prose 'Full Challenge suite proven on Genymotion' standing in for 9 named Challenges. This exactly the shape of the 2026-06-26 incid (.lava-ci-evidence/sixth-law-incidents/ 2026-06-26-c00-only-gate-shipped-broken flows.json) that 6.AK was written to preve recurring ONE CYCLE after the clause lan — because the gate that was supposed to it could not read the evidence format the used. Two further cycles, 1.3.17-1084 and 1.3.17-1085, declare NO commit SHA at a their evidence. NOT a live regression: it is historical finding about what was actually verified versus what was reported. Record the record of that release is honest. Sourc self-discovered during the LVA-149 blast-r assessment, 2026-08-26.
LVA-157	Bug	Queued	P0	WHAT: the three suites asserting 6.AA cla (test_phase_05_distribute_gate.sh, _vacuit _fix_audit.sh) each call write_evidence_re once and validate_evidence_record ZERO Until 2026-08-26, write_evidence_record stamped every record with anti_bluff_stat 'validated' — the SAME value the indepen validator writes on accept — so a record r examined was byte-identical on disk to on examined and accepted, and every consum read that field as a verdict. The aggregato compounded it by treating an absent statu 'not rejected', an asymmetry against 'resu where an unrecognised value already fell through to FAIL. NET EFFECT: clause 8 condition (B) — 'zero Evidence Records ca anti_bluff_status other than validated' — w

ATM ID	Type	Status	Severity	Description
				satisfiable by records that were never val INCLUDING all four mandatory real-devic challenge records for app-debug, app-rele api-app-debug and api-app-release, every them Challenge00CrashSurvival. Clause 8 rule authorising a distribution to proceed debug to release WITHOUT an operator p WHY THIS IS P0: that is the 2026-06-26 ' only gate shipped broken flows' forensic a (.lava-ci-evidence/sixth-law-incidents/) reproduced ONE LEVEL ABOVE where th sweep found it — not in a release's evider but in the test fixture that CERTIFIES the authorising unattended release. 6.Z claus states a cold-start check is the minimum a explicitly NOT sufficient. ATTRIBUTION SETTLED by controlled A/B, not inferred: swapping only evidence.sh + phase-02-tes gives pristine -> all three EXIT=0, fixed - three EXIT=1; the suites were relying on fail-open default. UNDERLYING FIX LAN the writer now emits a schema-legal fail-c placeholder and the aggregator classifies ways (validated / REJECTED* / anything e UNEVALUATED, which fails the phase). C the fixtures need repair, and it is not yet established whether clause 8 can still qua C00-only evidence once they validate prop — that is being determined now. Source: s discovered during the tag.sh critical fail-c fixes, 2026-08-26.
LVA-158	Bug	Queued	P1	WHAT: .gitmodules declares 'ignore = untracked' for submodules, so 'git status' parent does not report untracked files left inside a submodule working tree. The pip FR-000 precondition guard reads that sta decide the tree is clean, and therefore PA with untracked leftovers present in two submodules. STATUS CHANGE: the 2026 vacuous-pass sweep recorded this as P11 UNCONFIRMED; it is now CONFIRMED FIXTURE during the corpus-floor sweep. IT MATTERS: FR-000 exists so the pipelin never builds, tests or distributes from a tr carrying uncommitted state — a run that verifies artifacts built from unrecorded lo changes attributes its evidence to a comm does not describe what was actually built is the same attribution failure class as LV

ATM ID	Type	Status	Severity	Description
				(a run report claiming a commit_sha the r never tested), reached by a different rout Composes with the T054 re-review finding every push and both B-2 residuals in adva all-submodules.sh sit behind has_local_mo a precondition the script neither checks n states. Source: self-discovered, sweep find P11 promoted from UNCONFIRMED to CONFIRMED, 2026-08-26.
LVA-159	Bug	Queued	P1	WHAT: the test's 'git' stub does not honou -z flag, so the file list it feeds the three se 6.R scanners (scan-no-hardcoded-uuid.sh, no-hardcoded-ipv4.sh, scan-no-hardcoded hostport.sh) is not parsed as the scanners expect, and they examine ZERO files. The has therefore been green while asserting nothing about credential or hardcoded-lit detection. PROVED BY HEXDUMP during 2026-08-26 corpus-floor sweep; it became visible only because the new corpus floor a zero-file scan refuse instead of silently succeeding. SCOPE, stated precisely: the measured is in the TEST'S STUB, so what established is that the TEST was vacuous NOT that the production scanners are bro against a real corpus. Confirming product behaviour is separate work and is NOT cl here. WHY IT MATTERS: section 6.H trea credential leakage as a security incident a section 6.R forbids hardcoded literals; a s that certifies those scanners while feeding nothing is the highest-consequence instan the vacuous-pass family found this cycle. FIXED: the file is staged by another agent concurrent work; the remedy is one line i stub. Source: self-discovered during the c floor sweep, 2026-08-26.
LVA-160	Bug	Queued	P0	WHAT, MEASURED not argued (CASE 19 in tests/pipeline/ test_phase_05_distribute_gate.sh, re-meas every run rather than asserted in prose): 'records naming a Challenge OTHER than Challenge00CrashSurvival: 0', the gate re 'CONDITION (C) Device evidence for BOT variants : PASS' and 'clause 8 QUALIFIES 3)'. Condition (C) at scripts/pipeline/phase distribute.sh:908-1022 checks variant bin category, result, anti_bluff_status and the container-or-VM runner — it NEVER read

ATM ID	Type	Status	Severity	Description
				test_id. WHY THIS IS P0: clause 8 is the r authorising a distribution to proceed from debug to release WITHOUT an operator p and clause 8(C) states IN ITS OWN TEXT cold-start 'does not by itself satisfy 6.AK'. implementation does not enforce what the asserts. This is the 2026-06-26 'C00-only g shipped broken flows' incident condition, reachable through the gate written to pre it. AGGRAVATED BY THE 2026-08-26 AMENDMENT: the re-lettering withdrew condition (D) — the only condition gesturi per-variant channel verification — recordi as 6.AA-pipeline-debt-D OWED rather tha replacing it, on an analysis stating that '(C becomes MORE load-bearing, not less: wi retired it is the sole mechanical release-va guarantee'. That was true of the STRUCT and false of the SUBSTANCE, and the assumption was not tested before the amendment was applied. THE FIX (specif NOT implemented — the gate was deliber not touched): require, per variant, at leas PASSing real-device-challenge record who test_id is NOT the cold-start canary — equivalently, intersect the executed-PASS Challenge set with the cycle's claim-set pe 6.AK clause 1. Nothing implements this to Source: self-discovered during the clause-8 fixture repair, 2026-08-26.
LVA-161	Bug	Queued	P1	WHAT, measured live during the T062 end end run (run_id 2026-08-26T14-09-17Z): t test phase reported 74 real-device-challen FAILs. They are a TIMEOUT ARTEFACT, n broken features. phase-02-test-challenge. selects 73 Challenge classes in a SINGLE emulator-matrix invocation, and NEITHER NOR scripts/run-challenge-matrix.sh pass test-timeout (verified: zero occurrences in either file), so cmd/emulator-matrix's DEF 10 minutes applied. Attestation row: test_seconds=600.02, test_error='signal: Gradle's last log line: 'Tests 81/104 compl (6 skipped) (12 failed)'. Gradle was killed writing JUnit XML, so the phase parsed 0 and marked all 73 classes FAIL. Of what actually ran, 12 genuinely failed (e.g. Challenge46SearchTimeoutInterruptTest ComposeTimeoutException after 15000 m

ATM ID	Type	Status	Severity	Description
				UNCONFIRMED: the true status of the remaining ~61 — settled by raising --test- timeout or batching the classes. The api-a half is a SEPARATE, GENUINE failure (test_seconds=124.88, exit status 1), not timeout. WHY IT MATTERS beyond the w run: a future operator reading this eviden would conclude the product is broken wh not, and the same reasoning error in reve how a real breakage gets dismissed as 'ju timeout'. WHAT THE PIPELINE GOT RIGI worth preserving: it fabricated no passes killed classes, wording them 'No entry fou 0 file(s) parsed — either the class filter di match, or the run aborted before instrumentation completed'. Source: self- discovered during the T062 run, 2026-08-
LVA-162	Bug	Queued	P1	WHAT, measured during the T062 end-to- run: the per-AVD attestation row is marke gating: true but carries diag = {} and failure_summaries = []. Section 6.I.4 (Gro extension) requires every row to carry diag.target / diag.sdk / diag.device / diag.adb_devices_state — the per-AVD for snapshot captured immediately before instrumentation — plus parsed failure_summaries. scripts/tag.sh Group-E 3 exists to catch the AVD-SHADOW BLUF comparing diag.sdk against api_level and refusing when they differ; with diag empt comparison has nothing to read. The gate therefore inert on exactly the rows it is m police, and it is inert on a row that declar itself gating. This is the vacuous-pass fam gate reaching a verdict having examined items) landing inside the emulator attesta path, which the 2026-08-26 sweep did not because it swept scripts rather than emitt evidence. Composes with LVA-160 (clause condition (C) never reads WHICH Challen ran): two independent device-evidence ga each unable to see the field it exists to ch Source: self-discovered during the T062 r 2026-08-26.
LVA-163	Task	Queued	P1	WHAT, measured 2026-08-26 after initiali the previously-uninitialised nested submo at their already-recorded pins (zero pins moved): the recursive submodule set carr remotes, of which 41 are own-org (vasic-d

ATM ID	Type	Status	Severity	Description
				HelixDevelopment) and 28 are THIRD-PA They include github.com/google/perfetto, bytedance/UI-TARS and UI-TARS-desktop square/leakcanary, appium/appium, Genymobile/scrcpy, chroma-core/chroma, framework/allure2, anthropics/anthropic- quickstarts, SigNoz/signoz, Skyvern-AI/sk run-llama/llama_index, mem0ai/mem0, kiwitcms/Kiwi and 15 more - nearly all ve under submodules/helixqa/tools/opensour PLUS THE TOP-LEVEL SUBMODULE NA superspec, which points at github.com/ WangX0111/superspec. IMPACT: the oper directive 'commit and push all Submodule recursively to all upstreams', executed lit would attempt a push to 28 repositories o by other people and organisations. Such attempts would almost certainly fail for la write access, but attempting them is wro a credential that DID grant access would it an incident rather than a failed comma CONVERGING MEASUREMENT from the design pass: the 27 previously-uninitialise submodules are EXACTLY submodules/he tools/**, and every one has a third-party upstream - so they are precisely the 27 th CLAUDE.md's Automated Pipeline Pin-Ad Path condition (C) forbids advancing. Init them makes them RECORDABLE, never ADVANCEABLE. NOTE the governance d deny introduced by LVA-138 is what conta this today: superspec was one of the seve submodules named by NO operator decis and it points at a third-party repository - t allow-list defaulting to deny is doing real OPERATOR DECISION 2026-08-26: push org remotes ONLY, and REFUSE each thi party remote with a named per-remote re rather than skipping it silently - 'all subm handled' without per-submodule records i evidence, by the same reasoning section 6 applies to per-AVD attestation rows. Align section 6.W (which governs where WE pu and with condition (C). FILING NOTE: the version of this record was mangled by unescaped backticks in the authoring she command, which silently dropped the submodule name in two places. Corrected recorded because a tracker entry that qu

ATM ID	Type	Status	Severity	Description
				lost a fact is the same defect class this cy been evicting all day.
LVA-164	Bug	Queued	P1	WHAT, measured 2026-08-26 immediately initialising the 29 nested submodules at t already-recorded pins: submodules/helixq reports a permanent dirty state at the par level (M submodules/helixqa), caused by ending artifact three levels down. CHAIN traced end to end: helixqa -> tools/opensc docling (gitlink pin 4f37eeee == actual H 4f37eeee, so NO pin moved) -> tests/data sources/pftaps057006474.txt. That file sh with MIXED line endings (825 CRLF amo 1380 lines, 75716 bytes) while docling's own .gitattributes declares '* text=auto', normalises on checkout and the worktree never match the index. Git states the mechanism itself: 'CRLF will be replaced the next time Git touches it'. DECISIVE T stripping CR from both the index blob and worktree file yields byte-identical content the difference is line endings ONLY - ther content change and no pin move. NOTE o instruments: 'git diff --ignore-cr-at-eol' do NOT neutralise this and reports 1 differin because the endings are mixed rather tha uniformly CRLF-at-EOL. An earlier check that flag as the discriminator reached the conclusion; the correct discriminator is a stripped byte comparison. Recorded beca using the wrong instrument and believing answer is the defect class this cycle has b evicting. WHY THIS BLOCKS: submodule helixqa is the ONLY submodule the gover allow-list permits advancing (LVA-138 inv the list to default-deny with helixqa as the entry), and the T054 round-2 fix made an unclean submodule working tree a REFUS default, opt-in only via a flag the pipeline passes. So the one advanceable submodu now un-advanceable. It also fails the pipe FR-000 clean-tree precondition at the par level. Neither was true before the init. REPRODUCIBILITY: inherent to docling@4f37eeee and its .gitattributes; i reproduces on any fresh clone on a platfo where git normalises. It is not local dama must not be 'fixed' by committing the normalised file into a third-party submodu

ATM ID	Type	Status	Severity	Description
				do not own. CANDIDATE REMEDIES, nor chosen: (a) set core.autocrlf=false / core.e for that submodule and re-check-out so th worktree matches the index; (b) add a loc info/attributes marking that path binary s stops normalising; (c) teach the clean-tree precondition to treat an eol-only differenc clean - REJECTED as a default because it weakens the very check that makes the ac safe, and would need to be provably scope eol-only; (d) do not initialise docling at all which conflicts with the operator's all-59- recursive decision. Source: self-discovere while verifying the consequences of the recursive init, 2026-08-26.
LVA-165	Bug	Queued	P0	WHAT: during the T054 architectural ana an agent wrote a shell expansion of the fo DOLLAR-BRACE GITLAB_TOKEN COLON DASH NO BRACE, intending a presence c That operator expands to the variable's VA when the variable is set, so a real GitLab personal-access token was printed into th session transcript. CONTAINMENT, verifi measurement not assertion: 0 tracked file match a token pattern (git ls-files piped to grep), and 0 scratchpad artefacts contain unredacted token. The token did NOT rea THIS IS A REPEAT. CLAUDE.md's 67th section-6.L cycle records the IDENTICAL mistake with LAVA_FIREBASE_TOKEN, fi LVA-5, where the operator chose to accep risk rather than rotate. The same footgun now fired twice in one codebase, in two different agents, months apart. That is the signal worth acting on: it is not an individ lapse, it is a shell idiom that reads like a s presence-check and is not one. OPERATO ACTION REQUIRED per section 6.H claus rotate the GitLab token. REMEDY WORTH CONSIDERING beyond rotation: a presen check must never use the COLON-DASH operator. The safe forms are a test on the variable being non-empty, or printing only fixed literal such as SET or UNSET. A gre over tracked scripts and agent-authored commands for COLON-DASH applied to a name matching TOKEN, SECRET, KEY, PASSWORD or CREDENTIAL would catch class mechanically rather than relying on

ATM ID	Type	Status	Severity	Description
				nobody making the mistake a third time. Source: self-discovered and self-reported acting agent, 2026-08-26.
LVA-166	Bug	Queued	P0	WHAT, proven against the current file: the operating precondition stated in the scrip own header at lines 116-125 says that wit CLEAN submodule tree the run is 'all loca undoable, nothing leaving this machine'. T FALSE. PUBLISH_LOCAL_MODS is last consulted at line 1810; the step-5 verify command is invoked later and is NOT gat it, so the verify step executes on the DEFA FLAGLESS, CLEAN-TREE path with the operator's live credentials in reach. The a agent demonstrated the capability by performing an SSH-authenticated git pus an ordinary child process (target was a lo bare repo over ssh-localhost; no real upst was contacted). The operator's SSH key is unencrypted on disk. 25 own-org upstrea reachable. WHY THIS MATTERS MORE T ANY SINGLE ESCAPE: review rounds 4 a both APPROVED the script conditional on passing --publish-local-modifications, on t stated basis that the clean-tree path creat commit and issues no push and had 'survi two full adversarial passes'. That conclusi rests on a safety property the code does n implement. Five review rounds and eleven proven escapes were all reasoned about w a boundary that does not exist. ADDITION MEASURED FACTS from the same analys environment scrubbing does NOT close th 'env -i' with HOME redirected to an empty directory STILL pushed successfully, beca OpenSSH resolves the tilde from the pass entry rather than from HOME; (2) path m DOES work - bwrap with HOME and /run tmpfs made both push routes fail while le network intact; (3) isolation closes ZERO eleven escapes, proven rather than argue because a fully sandboxed verify step still poisoned the commit headers - the SCRIP publishes afterwards; what isolation remo review 5's strongest argument for the futi guarding. RESIDUAL, stated honestly: a sandbox still has write access to the tree builds, so it can lie about the build; cache poisoning across the boundary is real and

ATM ID	Type	Status	Severity	Description
				unaddressed since the gradle cache must writable and a later unsandboxed build ex what it finds. UNCONFIRMED: whether a Lava build completes inside the sandbox - experiment gates the whole item. Source: discovered during the T054 architectural analysis, 2026-08-26.
LVA-167	Bug	Queued	P1	WHAT: the credential-masking sandbox prototyped for LVA-166 is a GIT-CREDENT sandbox, not a credential sandbox, and th difference is load-bearing. Masking HOM run/user removes the SSH agent socket a on-disk key, which provably blocks git pus (probe fired 45s into a live 17-minute buil both routes exit 128, zero refs; falsified by unsandboxed control that succeeds, and b removing only the two tmpfs lines, after v both pushes land). But the build must be the repository tree, and that tree carries l secrets: .env (4072 bytes - Firebase token HMAC secret, tracker credentials), app/g services.json and api-app/google-services and keystores/release.keystore plus keyst debug.keystore. All are gitignored, so sec 6.H holds with respect to git - but gitigno does not withhold them from the build. PROVEN, not argued: the SANDBOXED b ran uploadCrashlyticsMappingFileReleas TWICE, i.e. it reached Firebase from insid sandbox. So the claim 'nothing leaving thi machine' remains FALSE even with crede masking in place. Whoever corrects the fa header recorded in LVA-166 must not sub a smaller falsehood. CONTEXT - the sand itself WORKS and this is not an argument against it: a real 4-variant Lava Android b completed inside it (BUILD SUCCESSFUL 17m 50s, exit 0, all four APKs apksigner-v with real Lava certs, 3-ABI Go+cgo liblava cross-compiled through the NDK), a cold- rebuild was byte-identical to the host's, an cache reuse survives (2682/3035 tasks UP DATE, ~54ms fixed entry cost, A/B wall-ti deltas of +42/+46/-42s on a ~130s build - flips, indistinguishable from noise). TWO FURTHER RESULTS worth keeping: (a) th architecture analysis's own prototype mas run WHOLESALE, but this repository live run/media/... - following it literally deletes

ATM ID	Type	Status	Severity	Description
LVA-168	Bug	Queued	P0	project and fails in 0.01s with 'cd: No suc or directory'. An implementer following it have recorded a FALSE NEGATIVE and discarded a working control. Mask /run/u never /run. (b) the cache-poisoning residu analysis disclosed and could not close IS closable: --ro-bind on the gradle cache fai Gradle WRAPPER needs a writable .lck ar build dies at 0s), but bwrap --overlay-src v tmp-overlay works - build exit 0, identical host cache verifiably unpoisoned, push sti blocked, at roughly 3x warm wall time. Th poisoning was confirmed end-to-end first: sandboxed process writes gradle init.d/ ar later UNSANDBOXED Gradle executes it the agent socket present. NOT ESTABLIS only the BUILD half was measured. phase test.sh under masking is UNCONFIRMED the higher-risk half, because section 6.AH emulators interact with the mount names Source: self-discovered during the LVA-16 sandbox feasibility experiment, 2026-08-2 WHAT, measured live during T062 attemp the pipeline run report never discloses a SKIPPED phase. A run invoked with --skip omits 'test' from phases[] ENTIRELY, so th report shows outcome PASS with every lis entry PASS - a run that executed ZERO te byte-indistinguishable in its report from o that ran the full suite and passed. Capture 2026-08-27T00-15-42Z, exit 0, preconditio PASS / build PASS / test SKIPPED / install PASS / live_verify PASS, with no 'test' entr phases[] at all. WHY THIS IS SERIOUS: s 6.AA clause 8 condition (B) requires that entry in phases[] is PASS and that zero Ev Records carry a non-validated anti_bluff_s A run that SKIPPED the phase which proc nearly all Evidence Records satisfies (B) vacuously - the array is all PASS because phase that could have failed is not in it. Th the vacuous-pass family, at the run-report inside the artifact that clause 8 reads to authorise an unattended debug-to-release distribute. It composes with LVA-160 (con (C) never reads WHICH Challenge ran) ar LVA-157 (clause 8's own certifying fixture never validated their records): three independent ways the same clause can qu

ATM ID	Type	Status	Severity	Description
				on evidence that proves nothing. THE FIX skipped phase MUST appear in phases[] w explicit SKIPPED result, and clause 8's evaluation must treat a SKIPPED phase a disqualifying rather than absent. An omitt phase and a passed phase must never be same bytes. SELF-CORRECTION recorde the measuring agent: it had earlier assert that a --skip run 'can never report outcom PASS', taken from a code comment rather from measurement. The measurement disproved the comment. Recorded becaus trusting a comment over an experiment is defect class this cycle has been evicting. Source: self-discovered during the T062 e end run, 2026-08-27.
LVA-169	Bug	Queued	P1	WHAT, measured during T062 attempt #4 scripts/pipeline/phase-02-test.sh writes la go/internal/*/.lava-ci-evidence/ directories are neither tracked nor gitignored. The pipeline's own FR-000 precondition requi clean working tree, so once the test phase run, EVERY subsequent run in that check refused - the pipeline dirties the tree it is to test and then cannot start again. This is identical shape as the pipeline-runs/ ignor defect closed earlier this cycle by a preve check-ignore guard, reached by a differen that guard probes .lava-ci-evidence/pipeli runs/ specifically and does not see these n per-package directories. Source: self-disc during the T062 end-to-end run, 2026-08-
LVA-170	Bug	Queued	P2	WHAT, measured during T062 attempt #4 LVA-162 fix added a Lava-side refusal for gating attestation row carrying an empty and it correctly FIRED in phase-02 on live phase-04 accepts the same empty-diag ro the same frozen submodule emitter (submodules/containers pkg/emulator ma writeAttestation). So one phase enforces s 6.I.4's diag requirement and another does on identical input. A gate that is enforced one path and not another is worse than on enforced nowhere, because its presence i coverage it does not provide. Source: self discovered during the T062 end-to-end ru 2026-08-27.
LVA-5	Bug		P0	

ATM ID	Type	Status	Severity	Description
		Operator-blocked		67th-cycle: the LAVA_FIREBASE_TOKEN default-expansion printed the token to the session transcript (NOT committed to git) Operator MUST rotate per §6.H clause 6 (firebase logout; firebase login:ci). **Sourc** discovered — .lava-ci-evidence/sixth-law-incidents/2026-05-20-firebase-token-echo leak.json OPERATOR DECISION, 2026-08 ACCEPTED RISK. The operator has revie this item and elected to ACCEPT the risk than rotate the credential at this time. Th is therefore NOT abandoned work and NO pending action awaiting an agent: it is a recorded risk-acceptance, retained in the register so it stays visible rather than bei closed and forgotten. WHAT IS AND IS N TRUE AS A RESULT: the leak itself remai matter of record — the token value was expanded into a session transcript on 2026-05-20 and was never committed to g the incident record cited above — and the decision does not retroactively make that exposure harmless. What the decision set the RESPONSE: no rotation is scheduled, no agent should treat the absence of rota an oversight or re-raise it as a new findin WHY THIS ITEM STAYS OPEN RATHER T CLOSING: closing it would assert the underlying condition was remediated, wh not what happened, and section 6.J forbid recording an outcome the work did not pr An accepted risk that has been closed is indistinguishable in the tracker from a ris was fixed, and that is precisely the confla this register exists to prevent. RE-OPENI CONDITION: if the accepted token is obs authenticating after the operator's intent changes, or if the operator later elects to this item resumes as ordinary pending wo under its original unblock options [A], [B] unchanged.